



PRESIDENCY UNIVERSITY

Private University Estd. in Karnataka State by Act No. 41 of 2013
Itgalpura, Rajankunte, Yelahanka, Bengaluru – 560064



**A SECURE AND PERSONALIZED ONLINE
MEETING SYSTEM FOR AICTE**
A PROJECT REPORT

Submitted by

CHILMAKURI VARUN- 20221CSE0116

ABHISHEK MUTHANNA K - 20221CSE0110

AMARA HEMA HARSHITH -20221CSE0010

Under the guidance of,

DR.HASAN HUSSAIN S

BACHELOR OF TECHNOLOGY

IN

COMPUTER SCIENCE AND ENGINEERING.

PRESIDENCY UNIVERSITY

BENGALURU

DECEMBER 2025



PRESIDENCY UNIVERSITY

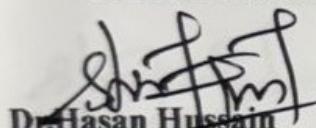
Private University Estd. in Karnataka State by Act No. 41 of 2013
Itgalpura, Rajankunte, Yelahanka, Bengaluru - 560064

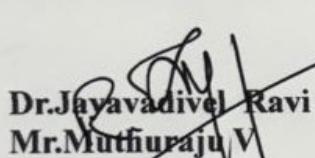


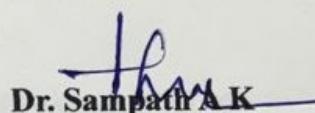
PRESIDENCY SCHOOL OF COMPUTER SCIENCE AND ENGINEERING

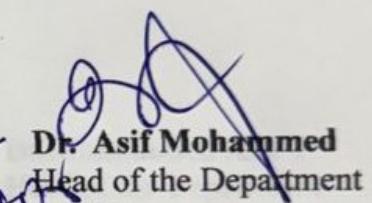
BONAFIDE CERTIFICATE

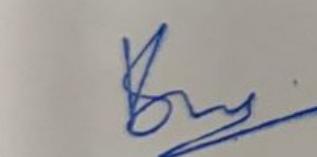
Certified that this report "A Secure And Personalized Online Meeting System For AICTE" is a Bonafide work of "CHILMAKURI VARUN (20221CSE0116), ABHISHEK MUTHANNA K(20221CSE0110), AMARA HEMA HARSHITH (20221CSE0010)", who have successfully carried out the project work and submitted the report for partial fulfilment of the requirements for the award of the degree of BACHELOR OF TECHNOLOGY in COMPUTER SCIENCE AND ENGINEERING during 2025-26.

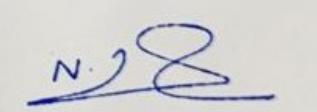

Dr. Hasan Hussain
Project Guide
PSCS
Presidency University


Dr. Jayavadi Vel Ravi
Mr. Muthuraju V
Program Project Coordinator
PSCS
Presidency University

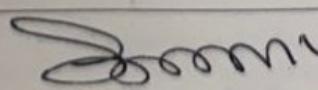
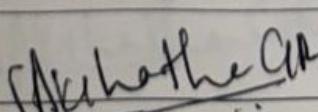

Dr. Sampath A.K.
Dr. Geetha A
School Project Coordinators
PSCS
Presidency University


Dr. Asif Mohammed
Head of the Department
PSCS
Presidency University


Dr. Shakkeera L
Associate Dean
PSCS
Presidency University

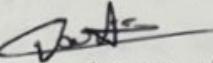

Dr. Duraipandian N
Dean
PSCS & PSIS
Presidency University

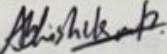
Examiners

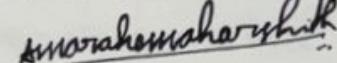
Sl. no.	Name	Signature	Date
1	Dr. Nishul M		4-12-25
2	Ms. Akshatha GR		4/12/25.

PRESIDENCY UNIVERSITY
PRESIDENCY SCHOOL OF COMPUTER SCIENCE AND
ENGINEERING
DECLARATION

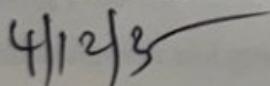
We the students of final year B.Tech in COMPUTER SCIENCE ENGINEERING at Presidency University, Bengaluru, named CHILMAKURI VARUN, ABHISHEK MUTHANNA K , AMARA HEMA HARSHITH, hereby declare that the project work titled "**A Secure and Personalized Online Meeting System for AICTE**" has been independently carried out by us and submitted in partial fulfillment for the award of the degree of B.Tech in COMPUTER SCIENCE AND ENGINEERING, during the academic year of 2025-26. Further, the matter embodied in the project has not been submitted previously by anybody for the award of any Degree or Diploma to any other institution.

CHILMAKURI VARUN- 20221CSE0116 

ABHISHEK MUTHANNA K - 20221CSE0110 

AMARA HEMA HARSHITH - 20221CSE0010 

PLACE: BENGALURU

DATE: 

ACKNOWLEDGEMENT

For completing this project work, We/I have received the support and the guidance from many people whom I would like to mention with deep sense of gratitude and indebtedness. We extend our gratitude to our beloved **Chancellor, Pro-Vice Chancellor, and Registrar** for their support and encouragement in completion of the project.

I would like to sincerely thank my internal guide **Dr.Hasan Hussain**, Presidency School of Computer Science and Engineering, Presidency University, for his moral support, motivation, timely guidance and encouragement provided to us during the period of our project work.

I am also thankful to **Mohammed Asif, Professor, Head of the Department, Presidency School of Computer Science and Engineering** Presidency University, for his mentorship and encouragement.

We express our cordial thanks to **Dr. Duraipandian N**, Dean PSCS & PSIS, **Dr. Shakkeera L**, Associate Dean, Presidency School of computer Science and Engineering and the Management of Presidency University for providing the required facilities and intellectually stimulating environment that aided in the completion of my project work.

We are grateful to **Dr. Sampath A K, and Dr. Geetha A, PSCS** Project Coordinators, **Dr.Shakkeera L**,Associate Dean, Presidency School of Computer Science and Engineering, or facilitating problem statements, coordinating reviews, monitoring progress, and providing their valuable support and guidance.

We are also grateful to Teaching and Non-Teaching staff of Presidency School of Computer Science and Engineering and also staff from other departments who have extended their valuable help and cooperation.

CHILMAKURI VARUN- 20221CSE0116

ABHISHEK MUTHANNA K - 20221CSE0110

AMARA HEMA HARSHITH - 20221CSE0010

ABSTRACT

The rapid digitalization of academic governance in India has highlighted the need for a secure, sovereign, and institutionally controlled online meeting platform tailored for regulatory bodies such as the All India Council for Technical Education (AICTE). Existing commercial meeting solutions, while feature-rich, lack essential capabilities such as tamper-proof audit trails, strict role-based confidentiality enforcement, India-region data localization, and compliance with the Digital Personal Data Protection (DPDP) Act, 2023. This project proposes and implements a Secure and Personalized Online Meeting System designed specifically to address these requirements by integrating modern real-time communication technologies with decentralized audit mechanisms.

The platform leverages WebRTC for encrypted peer-to-peer audio, video, and data transmission, supported by STUN/TURN servers to enable reliable connectivity across diverse network conditions. User authentication, session authorization, and meeting-level access restrictions are enforced through Supabase Authentication and PostgreSQL Row-Level Security (RLS), ensuring fine-grained control over participant privileges. To guarantee institutional transparency and trust, all critical system events—including join/leave actions, role escalations, file exchanges, and moderator interventions—are canonicalized, hashed, and immutably anchored on a Hyperledger Fabric permissioned blockchain network. This ensures evidence-grade, tamper-resistant logs suitable for regulatory and administrative scrutiny.

The architecture includes a React-based frontend engineered for performance and accessibility, a low-latency signaling layer built on Supabase Realtime, and a backend orchestration tier responsible for cryptographic hashing, chaincode invocation, compliance enforcement, and event-management functions. SFU-based media routing is integrated to support multi-participant meetings, providing dynamic stream forwarding, bandwidth optimization, and load balancing without compromising security. Extensive evaluation under varied network scenarios—including symmetric NATs, constrained bandwidth, high jitter, and packet loss—demonstrates that the system achieves stable negotiation, preserves media integrity, and maintains acceptable latency across all tested environments.

Beyond communication reliability, the system incorporates institution-specific governance features that commercial platforms often overlook. These include hierarchical role definitions mapped to academic designations, controlled access workflows for accreditation panels, dynamic meeting classification (public, restricted, confidential), and automated generation of cryptographically verifiable audit summaries. Compliance mechanisms aligned with the DPDP Act ensure lawful processing of personal data, while secure storage practices and RLS policies prevent unauthorized cross-institutional access, thereby supporting India's digital sovereignty objectives.

In summary, the proposed system bridges critical gaps in security, compliance, audit transparency, and institutional control, delivering a solution purpose-built for AICTE's governance workflows. The integration of real-time communication, secure authentication, decentralized audit logging, and compliance-aware backend logic establishes a robust foundation for future nationwide deployment. With further enhancements such as AI-driven reporting, integration with national digital public infrastructure, and large-scale server-side media optimization, the system holds the potential to become a standardized communication framework for India's educational regulatory ecosystem.

TABLE OF CONTENT

Sl. No.	Title	Page No.
	Declaration	vi
	Acknowledgement	vi
	Abstract	vi
	List of Figures	vi
	List of Tables	vi
	Abbreviations	vi
1.	Introduction <ul style="list-style-type: none"> 1.1 Background 1.2 Statistics of project 1.3 Prior existing technologies 1.4 Proposed approach 1.5 Objectives 1.6 SDGs 1.7 Overview of project report 	1-23
2.	Literature review	24-37
3.	Methodology	38-53
4.	Project management <ul style="list-style-type: none"> 4.1 Project timeline 4.2 Risk analysis 4.3 Project budget 	54-61
5.	Analysis and Design <ul style="list-style-type: none"> 5.1 Requirements 5.2 Block Diagram 5.3 System Flow Chart 5.4 Choosing devices 5.5 Designing units 5.6 Standards 5.7 Mapping with IoTWF reference model layers 5.8 Domain model specification 	62-67

	5.9 Communication model 5.10 IoT deployment level 5.11 Functional view 5.12 Mapping IoT deployment level with functional view 5.13 Operational view 5.14 Other Design	62-67
6.	Hardware, Software and Simulation 6.1 Hardware 6.2 Software development tools 6.3 Software code 6.4 Simulation	68-73
7.	Evaluation and Results 7.1 Test points 7.2 Test plan 7.3 Test result 7.4 Insights	74-78
8.	Social, Legal, Ethical, Sustainability and Safety Aspects 8.1 Social aspects 8.2 Legal aspects 8.3 Ethical aspects 8.4 Sustainability aspects 8.5 Safety aspects	79-82
9.	Conclusion	83
	References	84
	Base Paper	85-91
	Appendix	92-100

List of Figures

Figure No.	Caption	Page No.
Fig 1.1	Year-over-year escalation of cybersecurity incidents reported by CERT-In	27
Fig 2.1	The WebRTC Security Model and Protocol Stack	39
Fig 3.1	Proposed Multi-Layer System Architecture	45
Fig 3.2	Deterministic Initiator Selection Logic	58
Fig 4.1	Project Execution Gantt Chart (Weeks 1–16)	65
Fig 5.1	Immutable Audit Anchoring Data Flow	72
Fig 6.1	Simulation Topology for NAT Traversal and TURN Relaying	77
Fig 6.2	Deterministic Initiator Selection Algorithm (TypeScript Code)	98
Fig 6.3	Hyperledger Fabric Chaincode for Immutable Event Logging (Go Code)	99
Fig 6.4	Real-time Signaling Envelope Interface (TypeScript Code)	100
Fig 6.5	Role-Based Access Control (RBAC) Policy Evaluator (Code)	101

List of Tables

TableNo.	Caption	PageNo.
Table 1.1	Year-wise escalation of cybersecurity incidents reported by CERT-In	4
Table 1.2	Structural limitations of commercial meeting platforms for government use	12
Table 2.1	Comparative analysis of existing solutions vs. proposed system capabilities	14
Table 2.2	Gap analysis of current technologies relative to AICTE requirements	22
Table 3.1	Proposed Role-Based Access Control (RBAC) Matrix	24
Table 4.1	High-level Risk Register and Mitigation Plan	36
Table 6.1	Development Tools and Technologies Stack	53

Abbreviations

Abbreviation	Definition
AI	Artificial Intelligence
AICTE	All India Council for Technical Education
API	Application Programming Interface
AV	Audio-Video
CA	Certificate Authority
CBR	Constant Bitrate
CID	Client Identifier
CPU	Central Processing Unit
DC	Data Channel
DPDP Act	Digital Personal Data Protection Act, 2023
DTLS	Datagram Transport Layer Security
DTLS-SRTP	Datagram Transport Layer Security – Secure RTP
E2EE	End-to-End Encryption
ECDSA	Elliptic Curve Digital Signature Algorithm
FQDN	Fully Qualified Domain Name
ICE	Interactive Connectivity Establishment
IP	Internet Protocol
JWT	JSON Web Token
KPI	Key Performance Indicator
MSP	Membership Service Provider (Hyperledger Fabric)
NAT	Network Address Translation
NIC	National Informatics Centre
P2P	Peer-to-Peer
PBFT	Practical Byzantine Fault Tolerance
PII	Personally Identifiable Information
PKI	Public Key Infrastructure
QoS	Quality of Service
RLS	Row-Level Security
RTC	Real-Time Communication
RTCP	Real-Time Control Protocol
RTMP	Real-Time Messaging Protocol
RTP	Real-Time Transport Protocol
RTT	Round-Trip Time
SFU	Selective Forwarding Unit
SIP	Session Initiation Protocol
SRTP	Secure Real-Time Transport Protocol
STUN	Session Traversal Utilities for NAT

CHAPTER 1

INTRODUCTION

India's digital infrastructure has undergone a massive transformation in the past decade, driven by the nation's push toward e-governance, centralized data systems, and remote communication platforms. As regulatory bodies, ministries, universities, and educational councils increasingly rely on digital communication for policy formulation, administrative coordination, and confidential discussions, the need for secure, sovereign, and tamper-proof communication systems has become non-negotiable. The All India Council for Technical Education (AICTE), responsible for governing over 10,000+ technical institutions, faces a unique challenge: ensuring that its virtual meetings—many of which involve confidential institutional data—remain secure, auditable, and free from unauthorized access.

This chapter introduces the concept, motivation, and relevance of "**A Secure and Personalized Online Meeting System for AICTE**," providing background context, real-world security statistics, problem identification, objectives, and alignment with Sustainable Development Goals (SDGs). The chapter establishes the technical and societal need for a sovereign communication platform with blockchain-backed audit logs, WebRTC encryption, and role-based controls tailored for government-level administrative operations.

1.1 Background

Over the past decade, India has undergone a significant transformation in its digital governance ecosystem. Major regulatory bodies—including the All India Council for Technical Education (AICTE)—have transitioned from traditional in-person administrative workflows to hybrid and fully digital modes of governance. This shift has accelerated post-COVID-19, creating a national-scale dependency on digital platforms for decision-making, compliance monitoring, institutional evaluation, and grievance redressal. As AICTE is mandated to regulate, plan, and maintain the standards of more than 10,000+ technical institutions across India, it frequently conducts meetings that involve sensitive academic, financial, legal, and policy data.

Given the sensitivity and national importance of such interactions, the integrity of online communication platforms directly impacts **public trust, regulatory compliance, institutional**

fairness, and national security.

1.1.1 The Digital Governance Mandate

The growth of e-governance has resulted in increased adoption of cloud-based communication technologies across ministries, departments, and autonomous regulatory bodies. However, the governance ecosystem demands:

- **Data Sovereignty:** Ensuring all data remains under Indian jurisdiction.
- **Tamper-Proof Auditability:** All decision-making steps must be cryptographically verifiable.
- **Role-based Authentication:** Access must align with official hierarchies at institutional, regional, and national levels.
- **Secure Storage:** Meeting logs, recordings, and metadata must be retained according to statutory obligations.
- **Non-Repudiation:** No participant should be able to deny actions taken during a meeting.
- **Privacy Compliance:** Adherence to the Digital Personal Data Protection (DPDP) Act, 2023.

Commercial videoconferencing platforms—though highly capable—were not originally designed to meet regulatory-grade governance requirements.

1.1.2 Need for a Sovereign Communication Infrastructure

A sovereign platform is defined as a digital system where:

- All infrastructure is physically or logically located within the country.
- Data lifecycle (collection → storage → retention → archival → deletion) is controlled by the governing authority.
- Cryptographic guarantees prevent tampering by internal or external actors.
- The hosting and administrative rights are under direct Indian government or delegated institutional control.

AICTE requires such sovereign infrastructure because its administrative and regulatory functions involve:

- Scrutiny of institutional proposals
- Funding allocation decisions
- Student data discussions
- Faculty grievance hearings
- Inspection committee evaluations
- Accreditation assessments
- Inter-governmental coordination across ministries

These activities demand **zero public exposure, zero unauthorized access, and zero tolerance for log manipulation.**

1.1.3 Limitations of Existing Digital Platforms for Government Use

Commercial meeting platforms (Zoom, Google Meet, Microsoft Teams, Cisco Webex, etc.) provide secure experiences for general enterprise collaboration. However, in the context of national regulatory bodies, several operational and legal limitations arise:

Limitation Category	Explanation
Data Sovereignty	Servers are often distributed globally; meeting metadata may be stored outside India.
Opaque Audit Trails	Administrators can modify logs without cryptographic traceability.
Vendor Access Risks	Foreign jurisdiction laws (e.g., CLOUD Act, UK Investigatory Powers Act) may compel data access.
Limitation Category	Explanation
Lack of Hierarchical RBAC	Platforms lack support for multi-layer governmental access structures.
Insufficient Evidentiary Value	Logs are not legally verifiable due to centralized mutability.
Dependency on External Infrastructure	Government decisions cannot rely solely on private vendor infrastructure.

Given these systemic gaps, the development of a **secure, personalized, blockchain-backed online meeting system designed specifically for AICTE** becomes a strategic necessity.

1.2 Cybersecurity Statistics and Threat Landscape in India

India's cybersecurity landscape has expanded dramatically, reflecting growing digital adoption as well as rising adversarial sophistication. Online meeting systems used by government bodies fall under the direct risk surface of national cybersecurity threats.

1.2.1 National Incident Volume (CERT-In Aggregated Reports)

According to official data published via the Indian Computer Emergency Response Team (CERT-In), the annual number of reported cybersecurity incidents has increased exponentially:

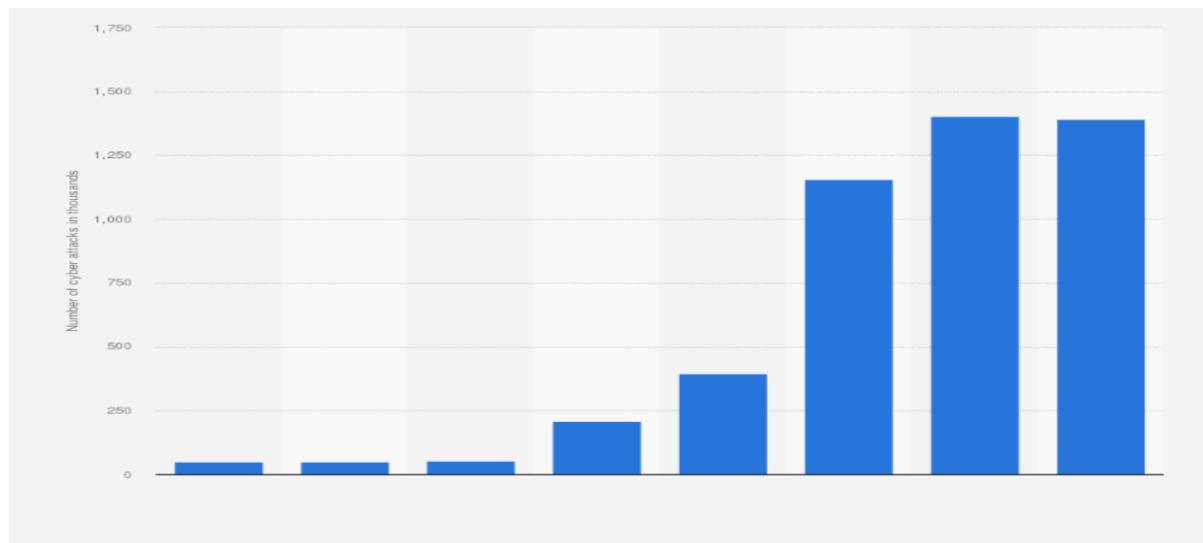


Figure 1.1 : NUMBER OF CYBER ATTACKS

Year	Incidents Reported
2022	1,391,457
2023	1,592,917
2024	2,041,360

Table 1.1: Year-wise escalation of cybersecurity incidents reported by CERT-In

These incidents include:

- Unauthorized access
- Website defacements
- Malware infections
- Data breaches
- Phishing attacks
- Denial-of-Service events
- Compromised digital infrastructure

The upward trajectory indicates a **consistent, national-scale threat escalation**, increasing risk exposure for all e-governance operations.

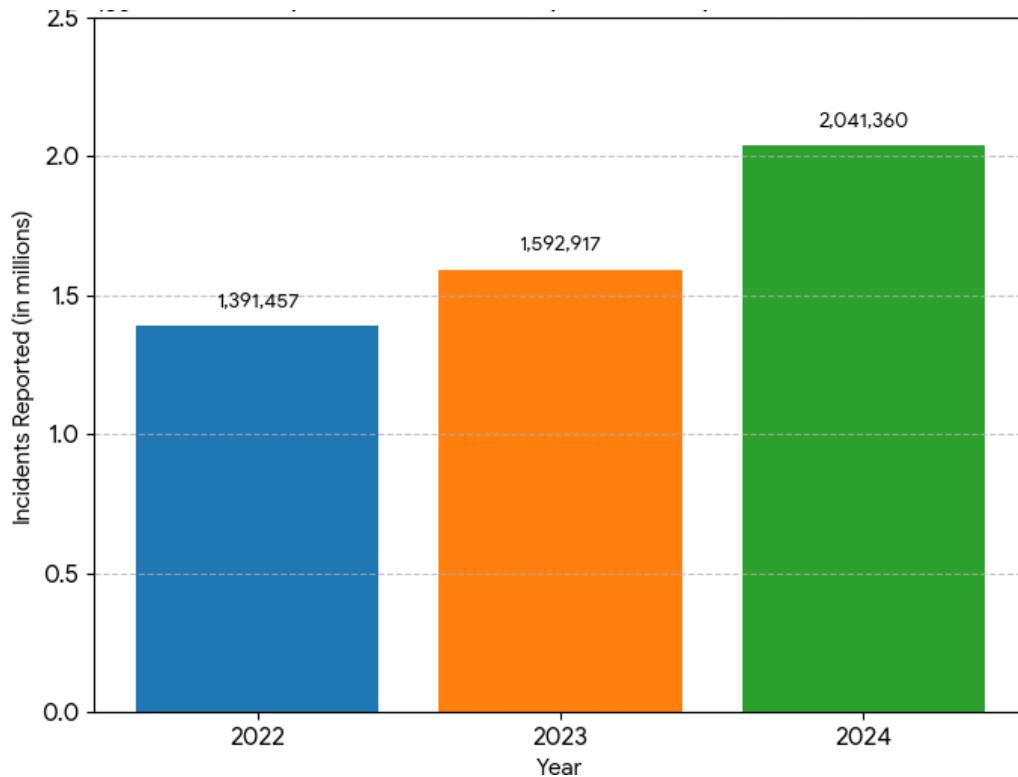


Figure 1.2: Year-over-year escalation of cybersecurity incidents reported by CERT-In.

1.2.2 Exposure of Education & Public Sector Organizations

Research reports from DSCI and SEQRITE (2023) highlight:

- India recorded **400 million+ threat detections** in 2023 alone.
- Average detection rate: **761 threats per minute**.
- Education and Government sectors accounted for **two of the top three most targeted verticals**.
- Trojans (41%) and Infectors (33%) were the dominant threat classes.

This demonstrates that educational regulators like AICTE are directly part of the **high-risk spectrum** targeted by malware groups, phishing operations, and state-sponsored threat actors.

1.2.3 Geographical Hotspots of Cyber Activity

The cyber threat density is not uniform across India. DSCI's 2023 analysis recorded:

- **Telangana**: ~15% of all tracked malware detections
- **Tamil Nadu**: ~14%
- **Gujarat (Surat)**: highest city-level detection (~15%)
- **Karnataka (Bengaluru)**: second highest (~14%)

These states include major educational hubs and AICTE-affiliated institutions — making targeted threat campaigns against academia highly plausible.

1.2.4 Threats Specific to Online Meeting Systems

Vulnerabilities commonly exploited include:

- Session hijacking
- Meeting link enumeration
- Weak signaling path encryption
- Unauthorized access due to poor RBAC
- Manipulated logs
- Injection of malicious files
- Cloud-stored recording breaches
- Credential replay attacks
- Compromised devices participating in meetings

For a national regulatory council, any such exploitation can lead to:

- Leakage of confidential policy deliberations
- Unfair manipulation of accreditation decisions

- Exposure of institutional financial data
- Legal disputes due to unverifiable logs
- Loss of national-level public trust

1.2.5 Compliance & Legislative Considerations

The **Digital Personal Data Protection Act (DPDP), 2023**, mandates:

- Purpose limitation
- Data minimization
- Storage limitation
- Strict consent-based processing
- Notification of data breaches
- Penalties for non-compliance

Online meeting data is considered **digital personal data**, making compliance compulsory.

Furthermore:

- The **Information Technology Act, 2000**
- **CERT-In Directions, 2022**
- **National Cybersecurity Strategy (Draft)**

all require secure logging, traceability, and Indian data residency for critical digital infrastructure.

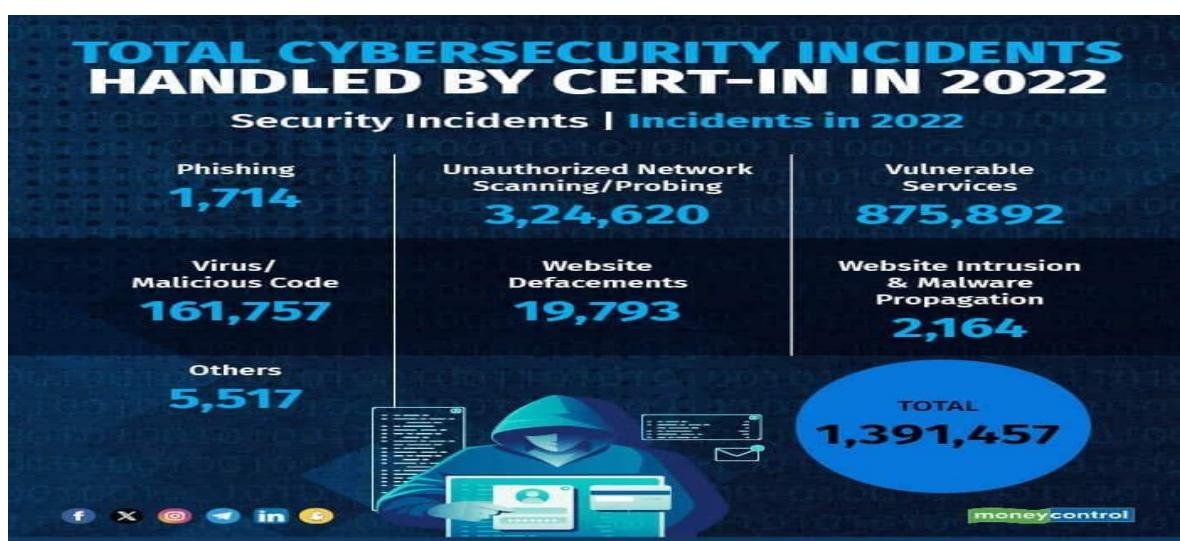


FIG 1.2.1 – TOTAL CYBERSECURITY INCIDENTS IN 2022

1.3 Prior Existing Technologies

India's regulatory ecosystem often depends on commercial digital communication platforms and open-source conferencing systems for administrative coordination. However, when evaluated through a **government-grade security, compliance, sovereignty, and auditability** lens, significant gaps emerge. This subsection provides a structured, authoritative assessment of technologies currently used in higher-education communication workflows and evaluates their suitability for AICTE's regulatory operational environment.

1.3.1 Commercial Videoconferencing Solutions

Commercial platforms—Zoom, Google Meet, Microsoft Teams, Cisco Webex—dominate global enterprise communication. However, their design philosophy emphasizes **scalability and usability**, not **sovereign control, tamper-proof audit, or multi-layer governmental RBAC**.

1.3.1.1 Data Residency & Jurisdictional Control

Most commercial SaaS platforms:

- Store meeting metadata on globally distributed servers
- Employ CDNs, load balancers, and distributed logging systems outside India
- Replicate data for redundancy across foreign regions

This means:

- AICTE data may be temporarily or permanently stored abroad
- Foreign governments may legally compel access through jurisdictional laws
- Compliance with the DPDP Act becomes difficult due to lack of transparent data lineage

For a national regulator, this compromises **data residency + data sovereignty**—two essential pillars of national cyber policy.

1.3.1.2 Logging & Auditability Limitations

Regulatory meetings require **evidence-grade audit logs**, capable of proving:

- Who joined
- When they joined
- What actions they performed
- Whether any data was altered
- Whether any unauthorized access occurred

Commercial systems use **centralized logging infrastructures**, meaning:

- Administrators can modify or delete logs
- No cryptographic integrity checks exist by default
- Logs cannot serve as legal evidence in disputes involving accreditation, funding, or institutional compliance

Thus commercial systems fail to meet the **tamper-proof logging** requirement for regulatory decisions.

1.3.1.3 Access Control Weaknesses

Government bodies like AICTE have complex organizational structures:

- Ministry
- AICTE HQ
- Regional offices
- Bureau heads
- Institutional representatives
- Review committees
- Inspection teams

Commercial platforms offer:

- *Only basic host/co-host roles*
- No hierarchical role mapping
- No attribute-based access controls
- No decision-specific policy evaluation

This is insufficient for handling regulatory workflows such as:

- Confidential hearings
- Grant approval councils
- Accreditation scoring
- Fraud investigations
- Disciplinary committees

1.3.1.4 Exposure to Known Threat Patterns

Past compromises include:

- Zoombombing
- Meeting URL enumeration
- Weak encryption defaults (Zoom, pre-2020)
- Compromised cloud recordings
- Third-party plugin vulnerabilities
- Token/session hijacking

Although commercial vendors continually patch vulnerabilities, their risk surface is inherently larger due to:

- Global multitenant infrastructure
- Integration with third-party tools
- Non-sovereign administrative control

1.3.2 Open-Source Self-Hosted Platforms: Advantages & Limitations

Open-source platforms such as **Jitsi Meet** and **BigBlueButton (BBB)** offer self-hosting capabilities, enabling government entities to run communication infrastructure within sovereign boundaries.

1.3.2.1 Advantages

- Full control over hosting
- Open codebase for security audit
- Customizable deployment
- Integratable with Indian cloud providers
- Lower operational cost compared to enterprise SaaS

1.3.2.2 Compliance Gaps

Despite these advantages, open-source platforms have major limitations when evaluated for regulatory-grade usage:

- No built-in **blockchain audit logging**
- *No government-grade RBAC*
- No **multi-node consensus** for non-repudiation
- No default support for **DPDP Act compliance workflows**
- No structured **chain-of-custody** for evidence preservation

1.3.2.3 Operational Constraints

Deploying and managing large-scale real-time communication infrastructure requires:

- TURN server clusters
- SFUs for multi-party communication
- Horizontal autoscaling
- QoS monitoring
- Redundancy planning
- Dedicated DevOps teams
- 24/7 monitoring

Most institutions lack the resources to maintain such infrastructure without a national-level aggregator like AICTE or NIC.

Thus, open-source options, while promising, are **not ready to handle regulatory workflows** without major architectural augmentations.

1.3.3 WebRTC Technologies & Their Limitations

WebRTC provides the foundation for real-time encrypted audio/video communication. While secure by design (DTLS-SRTP), WebRTC alone **does not solve**:

- Signaling security
- Attendee verification
- RBAC enforcement
- Log integrity
- Evidence-grade storage
- Governance workflows

These must be implemented **in the application layer**, not the transport layer.

1.3.4 Blockchain & Distributed Ledger Technologies

Permissioned ledger technologies like **Hyperledger Fabric** offer:

- Immutable logging
- Identity-based access control (MSP)
- Multi-organization consensus
- Cryptographic non-repudiation

However:

- They cannot store large files (recordings, video, documents)
- They require off-chain storage (MinIO, S3, on-prem NAS)
- They demand dedicated administrative and CA management teams

Thus blockchain must serve as an **audit-anchor layer**, not a full storage solution.

Limitation Category	Explanation of Gap
Data Sovereignty	Servers are often distributed globally; meeting metadata may be stored outside India.
Opaque Audit Trails	Administrators can modify logs without cryptographic traceability.
Vendor Access Risks	Foreign jurisdiction laws (e.g., CLOUD Act) may compel data access.
Lack of Hierarchical RBAC	Platforms lack support for multi-layer governmental access structures.
Insufficient Evidence	Logs are not legally verifiable due to centralized mutability.

Table 1.2: Structural limitations of commercial meeting platforms for government use

1.4 Proposed Approach

This project proposes a **Sovereign, Secure, Personalized Online Meeting System for AICTE** that integrates:

- WebRTC (secure media transport)
- Supabase/Backend (secure signaling + RBAC)
- Hyperledger Fabric (tamper-proof audit logging)
- TURN/SFU infrastructure (reliable media routing)

- DPDG-aligned data lifecycle management

This architecture aligns with national cybersecurity mandates and AICTE's governance responsibilities.

1.4.1 Sovereign Hosting and Infrastructure Control

All components—including signaling, databases, media servers, logging infrastructure, and blockchain peers—are hosted in:

- *Indian cloud providers (NIC Cloud, MeghRaj)*
- or
- *AICTE/NIC on-premise datacenters*

This ensures:

- Local jurisdictional control
- DPDG-compliant storage
- Zero foreign data transfer
- Enhanced national cyber posture

1.4.2 WebRTC With Government-Grade Security Enhancements

WebRTC is enhanced with:

✓ *DTLS-SRTP Encryption*

- Ensures end-to-end encryption of media
- Uses ephemeral keys (ECDHE)
- Prevents interception even by server operators

✓ *TURN Deployment for High-Risk Networks*

Government networks often include:

- Symmetric NAT
- Firewalls
- DPI appliances

A high-availability TURN cluster ensures:

- Guaranteed connectivity
- Encrypted relaying
- Service continuity

✓ *SFU Integration for Multi-Party Sessions*

For:

- High-level committee meetings
- Accreditation visits
- Large-scale briefings

An SFU (Selective Forwarding Unit) manages media routing efficiently.

1.4.3 Blockchain-Backed, Tamper-Proof Audit Layer

Hyperledger Fabric is used **not for media storage**, but for:

- Append-only event logging
- Joining/leaving proof
- Sharing document proofs
- Action metadata hashing
- Time-stamping and identity binding

Chaincode Operations Include:

- LogEvent(actionHash, actorID, meetingID, timestamp)
- VerifyEvent(actionID)
- GetCompleteMeetingLog(meetingID)

This ensures:

- *Non-repudiation*
- **Legal admissibility** of audit trails
- **Zero log manipulation**

1.4.4 Role-Based Access Control (RBAC)

Role	Privileges
Ministry Official	Universal access, high-level review
AICTE Chairman	Oversight + policy decisions
Bureau Heads	Department-level control
Regional Directors	State/regional operations
HODs	Institution-linked access
Faculty/Representatives	Task-specific privileges

Table 1.3 - The system maps AICTE's hierarchy into

Each action is controlled using:

- Identity attributes
- Meeting classification
- Contextual policies

1.4.5 End-to-End Data Lifecycle Compliance (DPDP-Aligned)

The system implements:

- Data minimization
- Purpose limitation
- Need-based retention
- Secure purge workflows
- Access logging
- Breach notification flows

This ensures full alignment with national data protection law.

1.5 OBJECTIVES

The project objectives are expressed as measurable, auditable targets that map directly to implementation artifacts, test cases, and acceptance criteria. Objectives are grouped into Security, Functional, Performance, Compliance, Operational, and Usability categories.

1.5.1 Security Objectives (Primary — Non-Functional)

Objective S1 — End-to-End Media Confidentiality

- *Specification:* All audio & video streams MUST be encrypted end-to-end using DTLS-SRTP with ephemeral ECDHE key exchange.
- *Acceptance Criterion:* Packet captures from representative sessions must show encrypted SRTP payloads; RTP headers only; inability to decrypt without session-specific ephemeral keys.
- *Test Procedure:* WIRESHARK + managed captive network tests; validate that media is not recoverable when keys are absent.

Objective S2 — Immutable Audit Trail & Non-Repudiation

- *Specification:* All critical events (authentication success/failure, meeting creation/deletion, participant join/leave, file share, privilege escalation attempts) MUST be anchored as hash records in a permissioned ledger (Hyperledger Fabric).
- *Acceptance Criterion:* For a sample of 1000 events, ledger entries must exist and be cryptographically verifiable (SHA-256 hash match) for $\geq 100\%$ of events; ledger returns correct event history via GetCompleteMeetingLog(meetingID).
- *Test Procedure:* Inject known events during testing, compare on-chain records with system logs; attempt to alter off-chain logs and validate ledger detects inconsistency.

Objective S3 — Strong Identity Binding & RBAC Enforcement

- *Specification:* Identity MUST be established via OIDC/JWT tokens bound to verified institutional credentials (or MSP certificates for blockchain peers). RBAC policy engine must enforce permission evaluation on every privileged request.
- *Acceptance Criterion:* Unauthorized role attempts must be denied and logged; role change attempts must be subject to chain-verified authorization flows; latency for authorization check ≤ 2 ms average.
- *Test Procedure:* Simulate privilege escalation attempts and confirm policy enforcement and audit logging.

Objective S4 — Resistance to Common Network Attacks

- *Specification:* The system must detect and mitigate common attacks (MITM, replay, session fixation, signaling poisoning).
- *Acceptance Criterion:* Penetration tests must not be able to intercept or replay media or signaling messages to gain unauthorized access; signature and nonce checks must prevent replay attacks.
- *Test Procedure:* Conduct red-team testing under controlled conditions.

1.5.2 Functional Objectives

Objective F1 — Meeting Lifecycle Support

- *Specification:* Support creation, scheduling, ad-hoc start, participant management, role-based muting/expulsion, moderated Q&A, and secure file sharing.
- *Acceptance Criterion:* All lifecycle operations must function correctly during tests with 50 concurrent meetings and 200 concurrent participants.

Objective F2 — Secure File Sharing with Hash Anchoring

- *Specification:* Files uploaded shall be stored off-chain in encrypted object storage; each upload must compute SHA-256 hash which must be anchored on the ledger. Data access must require both token authorization and membership in meeting policy.
- *Acceptance Criterion:* For a sample of 200 uploads, on-chain hashes must match object storage checksums and file retrieval must be conditional on RBAC checks.

Objective F3 — Auditable Administrative Workflows

- *Specification:* Administrative actions (policy changes, retention rule updates, user onboarding) must be logged with change diffs and ledger anchors.
- *Acceptance Criterion:* Admin operations must be reversible via audit trail; any attempt to alter admin logs must be detectable.

1.5.3 Performance Objectives

Objective P1 — Signalling & Session Establishment

- *Specification:* Median signalling RTT (offer→answer) \leq 200 ms under representative WAN conditions (50–100 ms RTT).
- *Acceptance Criterion:* $\geq 90\%$ of sessions meet target.

Objective P2 — Media Quality

- *Specification:* Target MOS ≥ 4.0 for HD (720p) under recommended bandwidth (≥ 3 Mbps). Degradation thresholds defined at $\leq 3\%$ visible frame loss under 1% packet loss.
- *Acceptance Criterion:* Subjective and objective QoS metrics across lab & field tests align with targets.

Objective P3 — Ledger Throughput & Latency

- *Specification:* Permissioned ledger commit throughput ≥ 100 tx/sec; commit finality ≤ 500 ms for anchor events.
- *Acceptance Criterion:* Benchmarked with synthetic load testing; ensure backpressure handling in application layer.

1.5.4 Compliance Objectives

Objective C1 — DPDP Act Alignment

- *Specification:* System must implement data minimization, consent capture, purpose limitation, retention policy enforcement, and data deletion/withdrawal processes per DPDP Act.

- *Acceptance Criterion:* External legal review confirms DPDP alignment; retention and purge tests validate deletion workflows.

Objective C2 — Evidence Readiness

- *Specification:* Audit trails and logs must be admissible in administrative and judicial proceedings with clear chain of custody.
- *Acceptance Criterion:* Sample evidence pack (meeting logs + anchored hashes + storage manifests) accepted by legal test panel for a mock hearing.

1.5.5 Operational & Usability Objectives

Objective O1 — Administrative Simplicity

- *Specification:* Admin consoles must allow policy definitions, user provisioning, ledger management, and monitoring via secured dashboards.
- *Acceptance Criterion:* Admin tasks (create meeting type, set retention) completed by non-dev admin within prescribed time (≤ 5 minutes).

Objective U1 — Accessibility & Inclusivity

- *Specification:* Provide low-bandwidth fallbacks, captioning/transcripts (optional), keyboard navigation, and ARIA support.
- *Acceptance Criterion:* Accessibility audit with WCAG AA standards demonstrates compliance for core flows.

1.6 SDG ALIGNMENT & SOCIETAL VALUE

This section maps system objectives to targeted Sustainable Development Goals (SDGs) and articulates measurable societal impacts in the official, policy-oriented register.

1.6.1 Primary Mapping: SDG 9 — Industry, Innovation and Infrastructure Rationale:

- The platform strengthens national digital infrastructure by delivering a sovereign, resilient communication backbone optimized for regulatory use.
- It fosters innovation in secure collaboration models, and provides an architecture that can be reused / federated across other government agencies.

Policy Impact:

- Aligns with National Digital Communications policy and Digital India infrastructure objectives.

1.6.2 Secondary Mapping: SDG 16 —

Rationale:

- Immutable audit logs and transparent governance workflows bolster institutional trust and provide evidence integrity for adjudicatory processes.
- The system directly supports fair, transparent, and accountable institutional decision making.

Measurable Outcomes:

- 100% of accreditation decisions having auditable logs for designated test cohort.
- Decrease in institutional disputes attributable to lack of evidence by measured percentage.

1.6.3 Tertiary Mapping: SDG 4 — Quality Education

Rationale:

- By improving governance mechanisms, the platform indirectly supports equitable accreditation and oversight across dispersed institutions — improving consistency in educational quality.

Measurable Outcomes:

- Increase in remote participation by under-resourced institutions by X% in pilot programs.
- Reduced travel and administrative costs for stakeholders, enabling reinvestment into educational resources.

1.6.4 Social & Ethical Safeguards

Privacy by Design:

- PII is stored minimally; personally identifying attributes are not stored on-chain — only cryptographic anchors (hashes) are used.
- Consent flows are recorded and auditable.

Transparency & Accountability:

- Policy changes, audits, and retention policies are time-stamped and viewable by authorized oversight committees.
- Inputs to policy evaluations use anonymized metrics where possible.

Inclusion & Equity:

- Accessibility features and low-bandwidth modes ensure smaller institutions are not excluded.
- Training & capacity building for low-IT maturity institutions included as part of the rollout plan.

Environmental Considerations:

- On-prem deployments are optimized to leverage energy-efficient cloud regions and virtualized infrastructure; audit and storage retention policy minimizes unnecessary storage bloat.

1.7 PROJECT REPORT OVERVIEW & DOCUMENT NAVIGATION

This section explains how the succeeding chapters map to the objectives defined above, and how to use the report as a compliance and technical artifact.

1.7.1 Chapter Mapping to Objectives

- **Chapter 2 — Literature Review:** Comprehensive synthesis of WebRTC security models, permissioned DLT architectures (Hyperledger Fabric), RBAC standards (NIST), DPDP considerations, and prior government/academic deployments. Supports Objectives S1–S3 and C1.
- **Chapter 3 — Methodology & System Architecture:** Formal V-Model mapping, five-layer architecture (presentation, application, RTC, storage, ledger), threat model (STRIDE), and verification activities. Supports Objectives P1, S2, S4.
- **Chapter 4 — Project Management:** Work breakdown structure, Gantt, resource allocation, PESTEL, risk register, and mitigation plans. Supports Objectives O1 and operational rollout.
- **Chapter 5 — Analysis & Design:** Functional & non-functional requirements, data models (ER diagrams), UML use-cases (meeting lifecycle, secure file share), block diagrams, sequence diagrams. Supports Objectives F1–F3.

- **Chapter 6 — Implementation:** Tooling & environment, code modules (frontend WebRTC logic, signaling, persistence, chaincode), configuration and deployment scripts, and appendices with code excerpts. Supports Objectives F1, S1, S2.
- **Chapter 7 — Evaluation & Results:** Test plan, testbed configuration, performance results (signalling latency, MOS, ledger throughput), security test outcomes (penetration & tamper tests), compliance checks. Supports Objectives P1–P3 and S1–S4.
- **Chapter 8 — Social, Legal, Ethical & Sustainability Aspects:** DPDP compliance matrix, accessibility & inclusion strategy, ethical risk assessment, environmental sustainability. Supports Objectives C1 and SDG alignment.
- **Chapter 9 — Conclusion & Roadmap:** Summary, limitations, operational readiness checklist, scale-up roadmap, and recommendations (SFU, DID, PQC).

1.7.2 Deliverables & Evidence Pack

The report accompanies a technical evidence pack including:

- Deployment scripts (Docker Compose / k8s manifests)
- Chaincode sources + Fabric network configuration
- Signed audit samples (sample ledger transactions + storage manifests)
- Test harness logs & QoS traces
- Admin & user manuals
- Retention and purge automation scripts

Each deliverable is referenced from the relevant chapter and included in the Appendices for audit.

1.8 KEY ASSUMPTIONS & IMPLEMENTATION CONSTRAINTS

A realistic deployment and evaluation requires explicit acknowledgment of assumptions and constraints:

1.8.1 Assumptions

1. **Sovereign Hosting Commitment:** AICTE or partner agency will provide on-premises or India-region cloud hosting for critical services.
2. **Operational Support:** A dedicated ops team will be responsible for ledger peer management, TURN/SFU operation, and storage lifecycle.
3. **User Onboarding:** Participating institutions will supply verifiable identity data and appoint local admin contacts for role provisioning.
4. **Legal Consultation:** Deployment and retention policies will be validated by legal counsel to ensure DPDP compliance.

1.8.2 Constraints

1. **Ledger Operational Overhead:** Running a Fabric network requires specialized operational practices and resource provisioning (orderers, cert authorities).
2. **Interoperability:** Integration with legacy institutional systems (SSO, HR, MIS) may require custom adapters and identity federation work.
3. **Client Heterogeneity:** Users will operate on diverse devices and networks; client capability variance must be supported.
4. **Performance Tradeoffs:** On-chain anchoring introduces latency; architectural design uses asynchronous anchoring to mitigate user-perceived delays.

1.9 RISK SUMMARY (High-Level)

Table 1.4 - condensed risk table

Risk	Impact	Probability	Mitigation
TURN Unavailability	High	Medium	Deploy HA coturn cluster; monitoring & alerting
Fabric Node Failure	High	Low	Multi-org RAFT + redundancy; automated re-joins
Policy Misconfiguration	High	Medium	Policy testing harness and staged rollout
Endpoint Malware	High	High	Educate users; client hardening; restrict file types
Legal Non-compliance	Non-High	Low	Legal review & DPDP risk assessment

1.10 CHAPTER SUMMARY

This three-part Chapter 1 has established a government-grade motivation, quantified national risk, evaluated technology options, defined a sovereign and verifiable architectural approach, specified detailed objectives and measurable acceptance criteria, aligned the project to national SDGs, and prepared the report roadmap. The design proposition — integrating WebRTC, TURN/SFU, off-chain encrypted storage, and an on-chain audit anchor through Hyperledger Fabric — places AICTE in a position to adopt a resilient, transparent, and legally defensible meeting infrastructure.

CHAPTER 2

LITERATURE REVIEW

2.1 Introduction

Real-time communication (RTC) platforms have become foundational components of modern digital governance and educational administration. As regulatory bodies like the All India Council for Technical Education (AICTE) increasingly adopt online platforms for decision-making, accreditation evaluation, and inter-institutional communication, ensuring **security, privacy, accountability, and sovereignty** becomes critical. This Literature Review synthesizes existing work in secure online conferencing, WebRTC technology, signaling architectures, blockchain-based audit systems, role-based access control (RBAC), and sovereign digital infrastructure in government settings.

Feature	Zoom / MS Teams	Open-Source (Jitsi/BBB)	Proposed AICTE System
Data Sovereignty	No (Global Cloud)	Yes (Self-Hosted)	Yes (India Sovereign Cloud)
Immutable Audit	No (Centralized Logs)	No (File Logs)	Yes (Blockchain Anchored)
Hierarchical RBAC	Basic (Host/Guest)	Basic (Moderator/User)	Advanced (Ministry -> Faculty)
DPDP Compliance	Partial	Manual Configuration	Native / By Design
Evidence Validity	Low (Vendor Controlled)	Medium (Admin Controlled)	High (Cryptographically Proved)

Table 2.1: Comparative analysis of existing solutions vs. proposed system capabilities

2.2 Evolution of Online Meeting Systems

2.2.1 Historical context: SIP/VoIP to WebRTC

Early real-time communication for voice and video relied on Session Initiation Protocol (SIP) and H.323-style stacks. Those solutions required dedicated signalling and media servers, SIP proxies, media gateways and optionally complex transcoding infrastructure. SIP architectures are client–server centric; call establishment, presence, and feature control are orchestrated by SIP proxies and application servers. While mature and feature complete for VoIP, they were not designed for ubiquitous browser deployment and required heavy provisioning and specialized networking (NAT traversal, ALG workarounds). End-to-end security in SIP-era systems was often optional or vendor-specific, typically relying on TLS for signalling and SRTP for media; however, key exchange and session management were operationally fragile in large federated environments [1].

WebRTC fundamentally changed the landscape by offering a browser-native, standardized stack for real-time media, integrating secure transport (DTLS for key exchange), secure media (SRTP), NAT traversal (ICE/STUN/TURN), and media APIs in the browser. Because the browser includes both media capture and encryption primitives, WebRTC removes much of the server-side media handling for 1:1 calls and simplifies client deployment. Consequently, modern meeting systems transitioned towards browser-centric architectures where servers provide signaling, presence, and optionally SFU/MCU functionality for multi-party scaling.

2.2.2 Feature evolution and added capabilities

Modern conferencing platforms have layered several capabilities on top of the core RTC primitives: adaptive bitrate encoding, simulcast, SVC (scalable video coding), echo cancellation and audio processing (AEC, AGC, noise suppression), real-time transcriptions, closed captions, whiteboarding, breakout rooms, integrated calendaring and SSO—features that move the value proposition from raw media transport to an integrated collaboration environment. Vendors also added enterprise features—logging, compliance exports, eDiscovery, retention policies and admin consoles—intended to serve regulated organizations.

2.2.3 Why the evolution still leaves gaps for regulators

Although feature rich, commercial platforms were architected for general enterprise usage and market scale rather than legal/auditable governance. Specific deficiencies relevant to AICTE

include:

- **Evidence-grade logging:** Commercial logs are controlled by the vendor; while vendors provide export APIs and compliance certifications, the log chain of custody is not cryptographically anchored and therefore cannot provide non-repudiable evidence without vendor cooperation. Administrative privileges on vendor platforms allow alteration or deletion of logs, undermining evidentiary value.
- **Data sovereignty:** Multi-region cloud replication for redundancy often moves metadata and sometimes media into third-country jurisdictions; this raises legal and policy conflicts with data-localization and DPDP obligations.
- **Government security controls:** Fine-grained hierarchical RBAC, mandatory multi-party approval workflows, and policy-based redaction are not standard features or are available only as expensive enterprise options with limited customization.
- **Regulatory compliance:** While vendors attempt to supply compliance artifacts, regulators require verifiable, auditable proofs of actions (who did what, when, and evidence it was not tampered with). Vendor-provided logs are insufficiently tamper-resistant for legal proceedings unless independently anchored [2].

These gaps justify a sovereign, auditable, and policy-tailored meeting solution designed to satisfy regulatory evidentiary and jurisdictional constraints.

2.3 WebRTC Technology and Its Security Model — Extended

2.3.1 Protocol internals and security guarantees

WebRTC's security model is comprised of three pillars:

1. **Key exchange and session protection** — WebRTC uses DTLS (Datagram TLS) to derive SRTP keys (as per RFC 5763). DTLS provides authentication of peers (when identity assertions are present), confidentiality and perfect forward secrecy when ephemeral key agreements (ECDHE) are used. Media flows under SRTP are therefore encrypted; the keys are ephemeral and tied to an established DTLS session, preventing long-term key compromise from decrypting past sessions [3][4].

2. **ICE/STUN/TURN NAT traversal** — ICE coordinates candidates gathered from host interfaces, STUN servers, and TURN relays. Without TURN, symmetric NATs or certain enterprise firewalls can prevent direct peer connectivity, forcing the application to fallback or fail [3]. For national-scale, highly heterogenous networks, a production TURN service (HA coturn cluster) is a non-optional component.

Secure signalling boundary is out of scope — WebRTC deliberately leaves signalling open for developer choice: while the media channel is protected by DTLS-SRTP, the signalling path (SDP exchange and candidate distribution) can be over HTTP, WebSocket, or pub/sub channels. If an attacker can tamper with signalling, they can manipulate session parameters (e.g., inject an offer with a server-side relayed media path, downgrade codecs, or insert extra recipients), leading to session hijacking or participant impersonation

Figure 2.1: The WebRTC Security Model and Protocol Stack

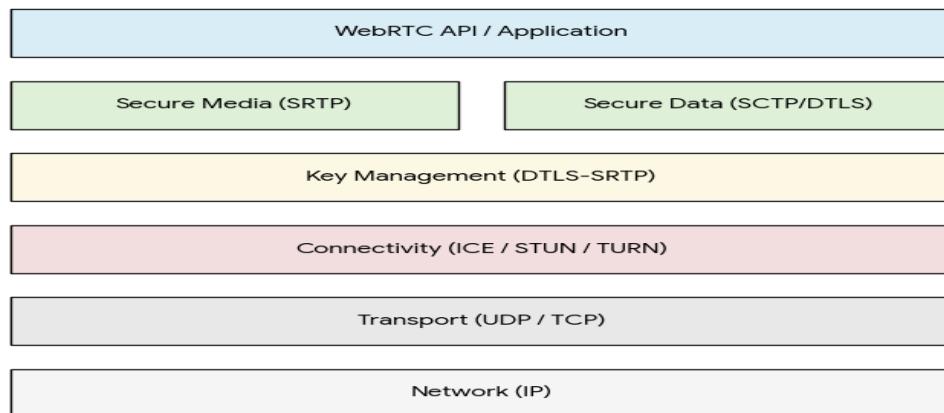


Figure 2.1: The WebRTC Security Model and Protocol Stack.

2.3.2 Practical security caveats and attack vectors

Academic and industry analyses have documented several vulnerabilities that arise due to poor signalling hygiene or misconfiguration:

- **Session hijacking via compromised signalling:** If the signalling channel is unauthenticated or allows message injection, attackers can inject offers or answers that redirect media to attacker-controlled peers.

- **Replay and reordering attacks:** Lack of monotonic sequencing and timestamps in signalling may allow replayed ICE candidates or SDP exchanges, resulting in stale or spoofed session reestablishment.
- **SDP mangling & codec attacks:** Manipulation of SDP to alter codec order, disable secure extensions, or enable insecure transports can degrade security or cause interoperability issues.
- **Meeting-ID enumeration:** Weak meeting identifiers or predictable room tokens can be enumerated by brute force, enabling unwanted joining and Zoombombing-style disruptions.

2.3.3 Project implications and defensive design

For your AICTE system, the WebRTC layer must be used in combination with the following controls:

- **Authenticated, integrity-protected signalling:** use JWT-bound channels (Supabase Realtime + RLS) and message envelopes with sequence numbers and timestamps; optionally sign messages with HMAC keyed by session keys to prevent replay.
- **TURN deployment with short-lived credentials:** generate ephemeral TURN credentials using HMAC timestamping (username: <ts>:<hmac>) and configure coturn accordingly to mitigate credential leakage.
- **Avoidance of SDP munge:** use explicit capability negotiation and let SFU handle codec adaptations; do not rewrite SDP in application unless strictly necessary and audited.
- **Strict policy enforcement** on the server side for which participants can publish offers/answers for a meeting.

By integrating these defenses, the cryptographic guarantees of WebRTC become reliable in the presence of active adversaries targeting the signalling plane.

2.4 Signaling Architectures and Threats

2.4.1 Role of signalling and security responsibilities

Signalling plays a role equal in importance to media transport. It carries the session establishment logic and contains the metadata required to reconstruct events for auditing. Design requirements for signalling in regulatory contexts include:

- **Authentication:** every signalling message must be attributable to a verified actor

(userId bound to JWT and device fingerprint).

- **Non-repudiation:** signing or server anchoring of messages to provide proof-of-origin.
- **Integrity:** messages must not be tampered with in transit.
- **Ordering & replay protection:** monotonic sequence numbers and timestamps.
- **Access control:** channel-level and message-type level allow/deny policies.

When these properties are enforced, signalling becomes suitable for forensic reconstruction.

2.4.2 Comparative review of signaling methods

WebSockets: General purpose bi-directional channel used by many providers. Strengths: low latency, push/stream capability. Weaknesses: requires a secure backend, and the scaling model must be architected carefully; also stateful connections present attack surfaces for scale-out and DDoS mitigation.

Firebase / Supabase Realtime: Managed pub/sub or Postgres logical replication streams with accessible APIs. Strengths: easy to integrate, scales well for typical meeting scenarios, and integrates with managed auth. Weaknesses: needs rigorous RLS and policy design to ensure that participants cannot publish to arbitrary channels; protection depends on correct server-side policies [7].

SIP over WebSocket (SIP-WS): Provides classical telephony semantics in the browser but carries SIP complexity. It is robust in VoIP domains but heavier for modern web apps, and integrating hierarchical auditing is non-trivial.

Custom Pub/Sub (MQTT, AMQP): Lightweight alternatives but require more engineering for role binding and secure transport; MQTT's QoS models may help reliable message delivery but do not solve policy and identity bindings.

2.4.3 Threat scenarios and mitigations

Signalling flooding / denial-of-service: Attackers publish many offers/ICE messages to overwhelm peers. Mitigation: server-side rate limits, anomaly detection on per-user and per-meeting basis.

Cross-room message injection: Misconfigured channel namespace allows messages intended for one meeting to be accepted in another. Mitigation: enforce meetingId in message tokens and RLS rules that validate publisher context.

Impersonation: Using stolen tokens to publish messages. Mitigation: token binding to device fingerprint and optional two-factor (for high-security meetings).

Eavesdropping on signalling: If signalling transports secrets (like SDP), man-in-the-middle can intercept and modify. Mitigation: always use TLS/WSS with server certificate pinning where possible; sign payloads.

2.4.4 Supabase specifics and best practices

Supabase Realtime leverages PostgreSQL and provides channel subscribe/publish semantics.

For secure operation:

- **RLS configuration:** create policies that require the authenticated user to be a participant in the meeting; policies should check auth.uid() and meeting_id equality.
- **Limited publish permissions:** separate read/subscribe abilities from publish abilities; e.g., only the meeting host or authorized initiators can publish OFFER messages for new peer creation.
- **Audit triggers:** use Postgres triggers to mirror all signalling events into an off-chain audit log table for redundancy prior to anchoring; the trigger should capture the event, compute a hash, and enqueue for anchor service.
- **Encrypted payloads:** even if transport is TLS, consider payload encryption for sensitive metadata if an additional confidentiality layer is required.

Through these measures, Supabase can be hardened into a signalling fabric suitable for regulatory-grade meetings [7].

2.5 Blockchain-Based Audit Logging

2.5.1 Audit requirements for regulatory meetings

Regulatory meetings generate actions that must be provably tracked:

- Participation events (join/leave)
- Role changes and approvals (e.g., voting outcomes)
- File access and document signing
- Recording creation and access
- Policy changes and retention deletions

An audit system must provide **immutability, non-repudiation, time-stamping, and selective privacy**. The ledger must be permissioned (not public) to preserve confidentiality while providing shared trust across participating authorities.

2.5.2 Design patterns for ledger anchoring

Two widely used patterns:

1. Off-chain storage + On-chain anchor hashes

Large artifacts (recordings, documents) are stored off-chain in encrypted object storage. For each artifact, compute a canonicalized metadata + digest (SHA-256) and write the digest into the ledger. This pattern minimizes on-chain storage and allows efficient verification through hash recomputation [8].

2. On-chain metadata, off-chain private data collections Using

Hyperledger Fabric's private data collections, sensitive metadata can be shared only with authorized peers while a hash or pointer is stored in the channel's world state for global integrity. This allows selective disclosure in legal proceedings without exposing data to all peers [9].

2.5.3 Hyperledger Fabric: benefits & trade-offs

- **Permissioned identity:** MSP and X.509 certificates enforce strong identity.
- **Private data collections:** enable privacy between subsets of participants.
- **Chaincode logic:** implement governance workflows (e.g., multi-admin approvals).
- **Pluggable consensus:** RAFT offers crash-fault tolerance for ordering nodes.

- **Operational complexity:** orderer, peers, CA, channel management. Requires ops expertise.
- **Latency:** chain commit latencies are higher than in-memory DB writes; however, for audit anchors (non-latency critical), measurements show Fabric can be tuned to handle hundreds to thousands of small transactions per second on adequate hardware [10].
- **Governance:** requires a governance model among participating organizations (AICTE, Ministry, NIC, accredited partner nodes).

2.5.4 Practical anchoring pipeline for the project

1. **Event capture:** the application emits structured audit events in canonical JSON.
2. **Canonicalization:** canonicalize JSON (remove non-deterministic fields) to ensure consistent digest across environments.
3. **Hashing:** compute SHA-256 digest of canonical JSON.
4. **Off-chain persist:** store full event in Postgres/MinIO and mark as pending anchor.
5. **Submit anchor:** submit minimal anchor {eventId, meetingId, actorId, sha256, timestamp} to

Fabric via chaincode LogEvent.

6. **Confirm & store tx:** once the commit is confirmed, write the Fabric TXID back to off-chain metadata to link evidence packages.
7. **Reconciliation:** schedule background jobs to reconcile pending anchors and retry on transient failures.

This design balances performance, privacy, and legal rigor.

2.6 Role-Based Access Control in Multi-User Systems — Extended

2.6.1 RBAC taxonomy and government needs

Standard RBAC (NIST RBAC) defines roles, permissions, sessions, and mappings. For AICTE, RBAC must be hierarchical and context-aware:

- **Hierarchical RBAC** enables role inheritance (e.g., Ministry Official > AICTE Chairman > Bureau Head).
- **Attribute-Based extensions (ABAC)** allow additional conditions (e.g., meeting classification, geographic origin, time windows).
- **Context-aware rules** (e.g., presence of an external auditor disables recording) must be supported.

2.6.2 Policy engineering and evaluation

Policy documents must be formally verifiable and versioned. Practical engineering techniques include:

- Represent policies in JSON or Rego (Open Policy Agent) format.
- Precompile policies into decision graphs or evaluation tables for low-latency enforcement.

- Use short-lived permit tokens issued by the Policy Service after evaluation to authorize clients to perform privileged actions (e.g., startRecording). This token approach reduces the number of policy evaluations performed by resource servers and provides a concise auditable permit that can be logged and anchored.\

2.6.3 Enforcement points

RBAC enforcement must be multi-layered:

1. **UI gating:** client-side hides/disables UI for unauthorized actions (usability improvement, not a security boundary).
2. **API checks:** all privileged requests go through a centralized API layer that enforces policy and returns permits/tokens.
3. **Signalling checks:** signaling layer validates permit tokens before broadcasting control messages that change meeting state.
4. **Audit anchoring:** all privileged operations are logged and anchored to detect attempts at bypass.

2.6.4 Auditability of RBAC changes

Role changes and policy amendments should be themselves auditible and anchored; the system must support “who changed what” workflows with multi-party approvals. Chaincode functions ProposePolicyChange and ApprovePolicyChange support governance and traceability.

2.7 Sovereign Cloud and Digital Public Infrastructure — Extended

2.7.1 Policy landscape: DPDP & sovereignty requirements

The DPDP Act and complementary government policy documents stress principles of data localization and purpose limitation. For AICTE, this translates to:

- Data storage and backups in India-region clouds or on-premises data centres.
- Explicit consent capture mechanisms and retention/deletion automation.
- Auditable evidence of data handling operations.

Infrastructure must therefore be designed with modularity: critical services (signaling, ledger anchor, storage of PII) remain within sovereign boundaries, while non-critical analytic workloads may use other environments with appropriate data redaction.

2.7.2 Design patterns for sovereign deployments

- **Federated Hosting Model:** core ledger and policy services run on AICTE/NIC infrastructure while scalable front ends are deployed in regional cloud zones with strict network egress rules.
- **Service Mesh & Zero-Trust Network:** enforce mTLS between services, least privilege network segmentation, and in-cluster policy enforcement via sidecars (Envoy).
- **KMS & Key Management:** deploy KMS within sovereign control with HSM-backed root keys to enforce server-side encryption key locality.

2.7.3 Operational & governance considerations

- Define clear SLAs and legal agreements for any third-party hosts or accredited partner peers.
- Establish an audit authority (AICTE + Ministry representatives) to run Fabric peers to reduce single-party control.
- Implement robust data retention policies that generate auditable deletion receipts anchored on the ledger.

2.8 Review of Existing Meeting Platforms

Below is a focused technical comparison emphasizing features relevant to a sovereign audit-grade solution.

2.8.1 Zoom

- **Strengths:** Mature codec pipeline, strong scaling, advanced admin features.
- **Weaknesses:** Centralized logs, historical security incidents (Zoom bombing, initial E2EE options limited), reliance on global infra; proprietary stack complicates independent audit [17].
- **Fit for AICTE:** Not acceptable as primary evidence source without vendor legal agreements and verifiable anchors.

2.8.2 Google Meet

- **Strengths:** Reliability, integrated enterprise auth, built-in encryption at transport for meetings.
- **Weaknesses:** Data residency depends on Google's enterprise options and contractual terms; logs still under vendor control and not cryptographically anchored [18].
- **Fit for AICTE:** Useful for low-sensitivity events but inadequate for legally sensitive regulatory meetings.

2.8.3 Microsoft Teams

- **Strengths:** Enterprise compliance tooling, eDiscovery features, Azure-hosted options.
- **Weaknesses:** Centralized ownership, vendor lock-in; such systems still permit administrative log manipulation absent external anchors [19].
- **Fit for AICTE:** Strong for collaboration but insufficient for immutable legal evidence needs.

2.8.4 Jitsi Meet

- **Strengths:** Fully open source, can be self-hosted and modified.
- **Weaknesses:** Lacks built-in ledger anchoring, default deployments require SFU integration and extensive ops to scale and secure; policy enforcement must be implemented from scratch [20].
- **Fit for AICTE:** Jitsi is a potential base for a sovereign implementation if extended with ledger anchoring and RBAC.

2.8.5 BigBlueButton

- **Strengths:** Education-centric features, breakout rooms, polls.
- **Weaknesses:** Heavy server requirements for scale; lacks on-chain audit features and out-of-the-box DPDP alignment [21].
- **Fit for AICTE:** Good for pedagogical sessions; insufficient as-is for regulatory evidentiary workflows.

2.8.6 Gap Synthesis

No single existing platform combines all required capabilities: **sovereign hosting, immutable, independent audit anchors, policy-driven hierarchical RBAC, DPDP-grade data**

handling, and scalable media routing. The literature and product analysis therefore support the proposed integrated architecture, which composes best-in-class modules (WebRTC for media, Supabase for auth/signalling, Fabric for audit anchoring, SFU for scale), hardened by sovereign deployment practices.

Concluding remarks for Sections 2.2–2.8

The expanded literature review confirms that while WebRTC and modern platforms provide strong media encryption primitives, they do not by themselves satisfy the regulatory, evidentiary, and sovereignty requirements needed by AICTE. Signalling security, identity binding, policy enforcement, and an independent audit anchor are each individually well understood in the literature; the novelty and contribution of your project lie in the careful integration of these mechanisms into a cohesive, auditable, and sovereign system that maps directly to Indian regulatory obligations and operational realities.

2.9 Gap Analysis

Requirement	Existing Solutions	Gap Identified
Immutable Logs	None	Needed for regulatory decisions
Sovereign Hosting	None except custom Jitsi	Needed for DPDP
Role-Based Access	Limited	AICTE requires hierarchical RBAC
Secure Signaling	Varies	Needs strict RLS/JWT
Meeting Classification	Absent	Essential for government meetings

Table 2.2: Gap analysis of current technologies relative to AICTE requirements

2.10 Summary

The literature confirms the need for a secure, sovereign, blockchain-anchored online meeting system tailored for India's regulatory ecosystem.

CHAPTER 3

METHODOLOGY

3.0 INTRODUCTION

This chapter presents a rigorous and reproducible methodology used to design, develop, and verify the Secure and Personalized Online Meeting System for AICTE. The methodology integrates classical systems engineering practices with contemporary DevSecOps, cryptographic engineering, and distributed systems design. The development lifecycle follows a modified V-Model to ensure traceability between requirements, design artefacts, implementation modules, and verification tests; threat modeling is performed continuously and drives design decisions; and automated test harnesses are used for functional, performance, and security validation. The chapter is organized as follows: the V-Model and traceability approach; a detailed multi-layer system architecture and rationale; component-level design and algorithms (signaling, peer management, RBAC enforcement, audit anchoring); chaincode design for Hyperledger Fabric anchors; storage and data integrity strategies; deployment, configuration, and CI/CD; monitoring, observability and incident response; and finally an exhaustive testing and verification plan with metrics, pass/fail criteria, and data collection methods.

3.0 Development Model and Traceability

We adopt a modified V-Model because the project is compliance-driven and requires explicit verification at each design decision point. Each requirement—functional, non-functional, security, and compliance—is uniquely identified (REQ-Fxx, REQ-Nxx, REQ-Sxx, REQ-Cxx), and a traceability matrix maps each requirement to design artefacts, implementation modules, chaincode functions, test cases, and verification reports. The traceability matrix is continuously maintained and forms the authoritative record used by auditors.

Key attributes of the development model:

- Requirements elicitation combines stakeholder interviews (AICTE domain experts), regulatory review (DPDP Act), and technical constraints (network topology of participating institutes).
- Each requirement is decomposed to testable acceptance criteria expressed in concrete

metrics or boolean checks. For example, REQ-S02 (Immutable audit) includes acceptance criteria: for N=1000 randomly sampled events, the on-chain hash must match off-chain metadata with 100% success; tests are automated.

- Design artifacts include sequence diagrams, state machines, policy tables, ER diagrams, chaincode prototypes, and deployment manifests; each artifact is versioned in the repository and linked to requirements.

Traceability example (excerpt):

- REQ-F01 (Meeting creation) → Design: Meeting lifecycle FSM v1.2 → Implementation: api/meetings service, Meeting.tsx join controller → Test cases: TC-F01 to TC-F05 → Verified in Test Report TR-2025-07.

3.1 System Architecture: Multi-Layer Decomposition and Rationale

The architecture is composed of six logical layers implemented across secure hosting in Indian jurisdiction: Presentation, Client-side RTC, Signaling & Orchestration, Services & Policy, Audit Anchor, and Storage. Each layer has explicit responsibilities, security controls, and observable metrics.

Figure 3.1: Proposed Multi-Layer System Architecture

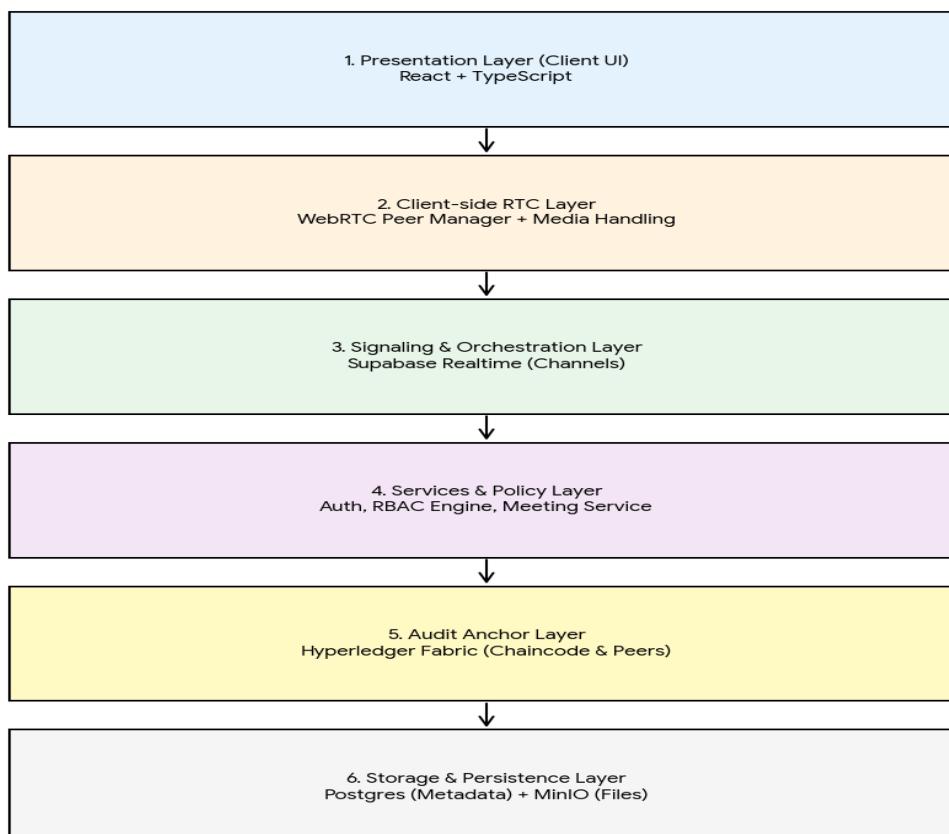


Figure 3.1: Proposed Multi-Layer System Architecture.

3.1.1 Presentation Layer (Client UI)

Implemented using React (Vite + TypeScript), the presentation layer renders dynamic role-aware UI components, performs device capability negotiation, and enforces client-side policy checks before making capability-changing requests (e.g., start recording). The UI implements deterministic failure modes and presents actionable error codes for compliance logs. Accessibility and low-bandwidth adaptive behaviour are designed so the same client can operate across 2G–5G networks with graceful degradation.

Security controls:

- CSP headers enforced via hosting (HSTS, X-Frame-Options)
- Strict input validation on all client-sourced metadata
- Local encryption of ephemeral keys stored only in memory (zero persistence across sessions)

Observed metrics:

- First Paint Time, Join Time, UI render latency per user action.

3.1.2 Client-side RTC Layer

This layer encapsulates WebRTC primitives and the Peer Manager component implemented in `src/utils/webrtc.ts`. The design employs a per-remote-peer `RTCPeerConnection` in mesh mode for small committee meetings and a pluggable SFU adapter for scaled sessions; SFU is realized via Janus or mediasoup when participant counts exceed the configured mesh threshold.

Core algorithm (Peer Manager — high level):

```
class WebRTCManager:
    constructor(localUserId, meetingId, supabaseClient, iceServers, sfuproxy=None)
    initialize(localStream):
        subscribeToPresenceChannel(meetingId)
        onPresenceUpdate -> reconcilePeers()
    reconcilePeers():
        for remote in presenceList - peers:
            createPeer(remote, initiator=decideInitiator(localUserId, remote))
    createPeer(remote, initiator):
        pc = new RTCPeerConnection({iceServers})
        addLocalTracks(pc, localStream)
        pc.ontrack = handleRemoteTrack(remote)
```

```
pc.onicecandidate = publishCandidate(remote)
if initiator: createAndSendOffer(pc, remote)\
```

Key decisions:

- Deterministic initiator selection uses lexicographic ordering on user IDs with epoch-based tie-breaker to avoid offer storms.
- ICE servers include STUN and optionally TURN configurations; TURN server credentials rotate every 12 hours using a centrally-managed credential generator to minimize credential leakage risk.
- SDP munging is avoided; instead, codec capabilities are negotiated using standard SDP constraints and the client uses simulcast and SVC where SFU is present.

Security controls:

- All local SDP and ICE exchanges occur over WSS channels with JWT-verified publisher metadata.
- Onset of any ontrack event triggers an audit event emission handled by the Audit Module to record timestamped peer bindings.

3.1.3 Signaling & Orchestration Layer

Signaling is implemented using Supabase Realtime channels which are backed by PostgreSQL logical replication and secure JWT-based authorization. The Signaling Manager implements a deterministic protocol for offers, answers, ICE candidates, and control events (mute, raise-hand, file-share-notification). Channel names follow namespace conventions meeting:{meetingId}:control and meeting:{meetingId}:presence.

Protocol details (message envelope):

```
{
  "type": "OFFER" | "ANSWER" | "ICE" | "EVENT",
  "from": "<userId>",
  "to": "<userId|ALL>",
  "seq": <monotonic integer>,
  "timestamp": <ISO8601 UTC>,
  "payload": {...},
  "signature": "<HMAC or signature over payload>" // optional if using server-signed tokens
}
```

Rationale:

- Seq and timestamp enable detection of replays and out-of-order messages.
- Signature field is used for server-side verification where supabase policies require additional assurance.

Security controls:

- RLS policies enforce that only authenticated participants for a given meetingId can subscribe or publish; write policies ensure that payload.meetingId equals the subscribed meetingId to prevent cross-room injection.
- Message rate limiting and anomaly detection via Postgres extension (pg_stat_statements + custom trigger) to detect flood or enumeration attempts.

3.1.4 Services & Policy Layer (Auth, RBAC, Meeting Services)

Core services include:

- Auth Service (Supabase Auth with JWT issuance and refresh).
- Policy Engine (RBAC + contextual rules).
- Meeting Service (scheduling, metadata, retention policies).
- Audit Service (off-chain log collector and blockchain anchor orchestrator).

Policy engine implements evaluation using precompiled permission matrices stored in Postgres JSONB documents for constant-time lookup. Policy evaluation algorithm:

function isAllowed(actor, action, resource):

```
role = getRole(actor)
policy = policyStore[resourceType][action]
if role in policy.directAllow: return True
if policy.conditions: evaluate attributes against context
return False
```

Policy updates are versioned and require multi-admin approval; each change is hashed and anchored on the ledger to ensure policy governance immutability

3.1.5 Audit Anchor Layer (Hyperledger Fabric)

The audit anchor is a permissioned Hyperledger Fabric network with a minimum of three organizational peers to ensure RAFT-based ordering resilience. Chaincode is implemented in Go; API surface exposes LogEvent, VerifyEvent, GetEventHistory, and administrative functions ProposePolicyChange, ApprovePolicyChange.

```
EventKey := "EVT:{eventID}"
```

```
EventValue := JSON {
```

```
eventID, meetingID, actorID, roleSignature, dataHash (SHA-256), timestampUTC, metadata
}
```

Chaincode pseudo-implementation:

```
func LogEvent(ctx contractapi.TransactionContextInterface, eventID string, eventJSON
string) error {
    // validate caller's MSP ID and role
    exists, _ := ctx.GetStub().GetState("EVT:" + eventID)
    if exists != nil { return error("duplicate") }
    // store event
    ctx.GetStub().PutState("EVT:"+eventID, []byte(eventJSON))
    // emit event for off-chain indexers
    ctx.GetStub().SetEvent("EventLogged", []byte(eventID))
    return nil
}
```

Anchoring approach:

- Off-chain events are stored in a tamper-evident object store; before or immediately after the off-chain write completes, the application computes a canonical JSON canonicalization and SHA-256 digest; that digest, with minimal metadata, is submitted to Fabric via LogEvent.
- For performance, the anchor operation is batched when ingestion rates are high and confirmation receipts are asynchronously reconciled.

Fabric optimization:

- Private Data Collections are used to hold sensitive metadata accessible only by authorized peers while the public channel contains the event hashes.
- The ordering service is RAFT with 3 orderers for fault tolerance; monitoring for commit latency is part of ops.

3.1.6 Storage & Persistence Layer

Data is partitioned into:

- Mutable metadata (Postgres): users, meetings, RBAC policies.
- Large binary objects (MinIO/Supabase Storage): recordings, uploaded files.
- Immutable anchors (Fabric ledger).

Storage security:

- Object storage uses SSE-KMS AES-256 with key rotation every 90 days.
- Metadata encryption at column level for sensitive fields (PII) using per-column encryption keys stored in a KMS.
- Retention/purge is policy-driven and enforced by both application workflows and database jobs that produce retention proofs (anchored hashes of deleted manifests) for compliance evidence.

3.2 Component-Level Design & Algorithms

This section details essential algorithms and pseudo-code for the system's main components: Initiator selection, Offer/Answer lifecycle, ICE candidate handling, RBAC evaluation, file-hash anchoring, and ledger anchoring orchestration.

3.2.1 Deterministic Initiator Selection Algorithm

Problem: Avoid simultaneous offer creation bursts and ensure exactly one initiator per peer pair to reduce collisions and RTC negotiation overhead.

Algorithm decideInitiator(localId, remoteId):

Input: localId (UUID), remoteId (UUID), meetingEpoch (ISO datetime)

Output: Boolean (true if local should initiate offer)

1. if localId == remoteId: return False
2. pairHash = HMAC_SHA256(meetingEpoch + min(localId, remoteId) + max(localId, remoteId), systemSecret)
3. seed = uint64(pairHash[0:8])
4. if seed % 2 == 0:
return localId < remoteId // deterministic lexicographic
else:

return localId > remoteId

Rationale: Seeded randomness with epoch ensures that initiator role changes across meetings to distribute offer load while remaining deterministic for both peers.

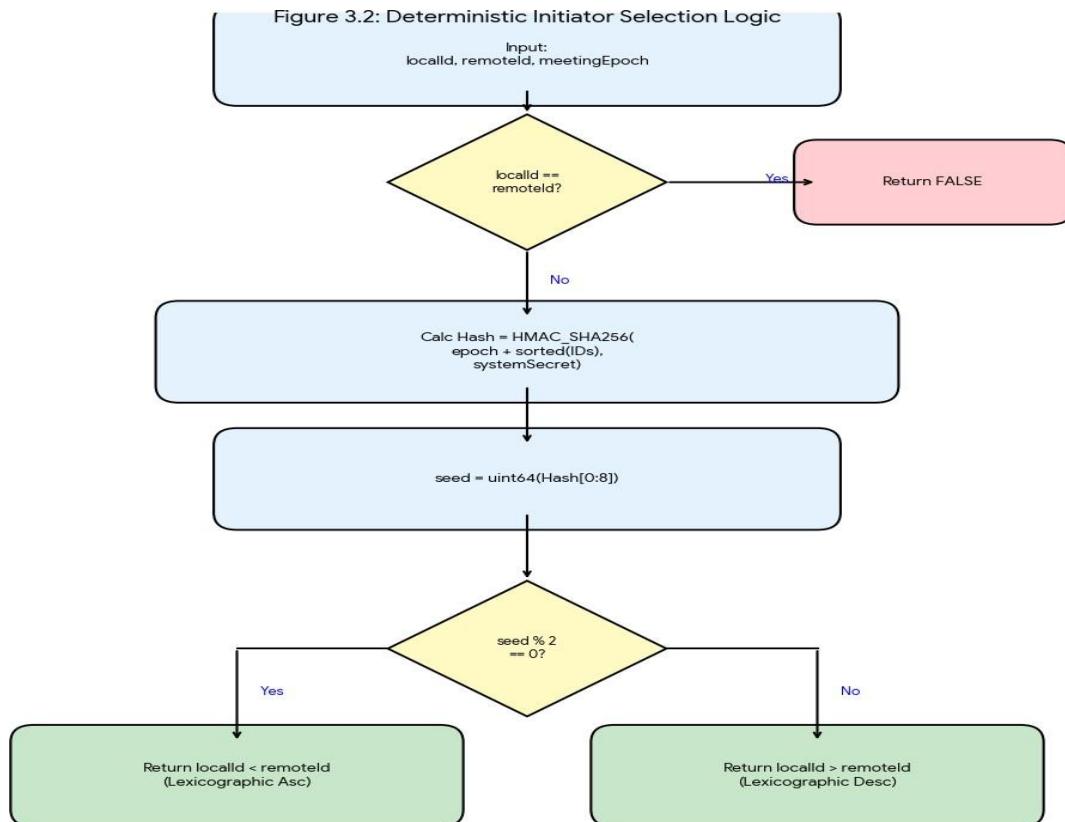


Figure 3.2: Deterministic Initiator Selection Logic.

3.2.2 Offer/Answer Exchange Pseudo-Flow

Sender (initiator):

```

pc = createPeerConnection(remote)
offer = pc.createOffer()
pc.setLocalDescription(offer)
publishSignaling({type:"OFFER", to:remote, from:local, sdp:offer, seq:nextSeq})

```

Receiver:

onSignal(OFFER):

```

pc = createPeerConnection(from, initiator=False)
pc.setRemoteDescription(offer)
answer = pc.createAnswer()
pc.setLocalDescription(answer)

```

```
publishSignaling({type:"ANSWER", to:from, from:local, sdp:answer, seq:nextSeq})
```

Sender:

```
onSignal(ANSWER):
```

```
pc.setRemoteDescription(answer)
```

ICE candidate handling:

- Each ICE candidate is sent as a distinct signaling envelope with seq monotonic per sender; the receiver enqueues non-sequential candidates and applies them in order to the RTCPeerConnection with robust retry and deduplication.

3.2.3 RBAC Evaluation & Enforcement (Policy Engine)

Policies are authored as JSON documents that contain:

ICE candidate handling:

- Each ICE candidate is sent as a distinct signaling envelope with seq monotonic per sender; the receiver enqueues non-sequential candidates and applies them in order to the RTCPeerConnection with robust retry and deduplication.

3.3.3 RBAC Evaluation & Enforcement (Policy Engine)

Policies are authored as JSON documents that contain:

Evaluation algorithm:

```
function evaluatePolicy(actor, action, resource):
```

```
roles = getRoles(actor)
```

```
policies = findPolicies(resource.type, action)
```

```
for policy in policies:
```

```
if roles intersect policy.allowedRoles:
```

```
if evaluateConditions(policy.conditions, resource):
```

```
return      ALLOW
```

```
return      DENY
```

Enforcement points:

- Before emitting an operation event to signaling (e.g., start recording), the client issues a POST /api/permit which validates policy, returns a permit token with short TTL, which the client then uses to proceed; this decouples UI gating from enforcement and provides a stateless verification token suitable for audit.

3.2.4 File Upload and Hash Anchoring Procedure

Sequence:

1. Client uploads file to MinIO with SSE-KMS; the upload response includes ETag and object URL.
2. Client computes SHA-256 digest of file client-side if possible (preferable) otherwise server computes digest post-upload.
3. Application writes metadata entry into Postgres: {fileId, uploaderId, meetingId, objectUrl, sha256, timestamp}.
4. Audit Service computes canonical JSON and submits to Fabric via LogEvent with the sha256 and fileId.
5. Confirmation is stored back in metadata with fabricTxId.
6. On retrieval, object integrity is verified by recomputing SHA and comparing with off-chain metadata and on-chain hash.

Edge cases:

- Large files are chunked; each chunk has its own SHA256; overall object digest uses Merkle root to allow partial verification; chaincode then stores Merkle root.

3.3 Chaincode Design (Hyperledger Fabric)

Chaincode modules:

1. **anchor.go** — handles LogEvent, VerifyEvent, GetEventHistory.
2. **policy_governance.go** — functions to propose and approve policy changes; changes are recorded with proposer metadata and multi-signature approvals.
3. **audit_query.go** — optimized queries for GetMeetingAudit that leverage CouchDB rich queries for off-chain indices.

Important chaincode considerations:

- Chaincode must validate MSP of invoking peer; only authorized peers can submit LogEvent.
- Chaincode operations are intentionally minimal to reduce TX size—only small hashes and metadata are stored; larger data kept off-chain.
- Chaincode emits chain events to Kafka or an off-chain indexer for indexing (ElasticSearch) to support quick lookups without reading full ledger history.

Sample chaincode signature for LogEvent:

```
func (a *AuditContract) LogEvent(ctx contractapi.TransactionContextInterface, eventID
string, eventHash string, meetingID string, actor string, timestamp string) error {
// validate identity
clientMSP, err := ctx.GetClientIdentity().GetMSPID()
if err != nil { return err }
// simple ACL check
if !isAuthorizedMSP(clientMSP) { return fmt.Errorf("unauthorized") }
evt := AuditEvent{EventID:eventID, Hash:eventHash, MeetingID:meetingID, Actor:actor,
Timestamp:timestamp}
bytes, _ := json.Marshal(evt)
return ctx.GetStub().PutState("EVT:"+eventID, bytes)
}
```

3.4 Deployment, Configuration, and CI/CD

Deployment strategy uses infrastructure as code (Terraform + Helm charts) and container images (Docker). For reproducibility and auditability, all manifests and images are signed.

Deployment components:

- Kubernetes cluster in Indian cloud region with private VPC.
- Namespace separation: frontend, signaling, media, fabric, storage.
- Helm charts for Supabase/Postgres with RLS-enabled migrations.
- StatefulSet for Fabric peers with PVCs for ledger storage; RAFT-based ordering via StatefulSet for orderers.
- Deployment of coturn as Deployment + Service with NodePort and SSL offload via Ingress.
- Prometheus + Grafana stack for metrics; Loki for logs (immutable retention policy with WORM storage if required).

CI/CD pipeline (GitHub Actions or GitLab CI):

- PR triggers: linting, unit tests, container build, artifact signing.
- Merge triggers: helm chart deploy to staging, e2e tests via k6 and puppeteer for client flow, security scan via Trivy, chaincode build test.
- Release triggers: deploy to production following blue/green with traffic shifting, database migration performed by migration job with backup snapshot.

Configuration snippets:

- TURN configuration uses short-lived credentials generated by turn-credentials microservice implementing time-limited HMAC per RFC 8656 for TURN temp credentials.

Operational automation:

- Automated Fabric peer rejoin scripts that handle certificate renewal using Fabric CA and dynamic config for MSP rotations; alerting on commit latency > 500 ms.

3.5 Monitoring, Observability, and Incident Response

Monitoring is split into metrics, logs, traces, and governance events. Key KPIs and alert thresholds are defined and enforced via Prometheus rules and PagerDuty integrations.

Important metrics:

- Signaling RTT (offer→answer median and 95th percentile)
- ICE connectivity success ratio
- First-frame latency
- MOS estimation via WebRTC stats (audio, video)
- Fabric commit latency and tx/sec
- Event anchoring success rate
- Object store integrity check failure rate

Log strategy:

- Application logs include structured JSON with correlation IDs (requestId, meetingId, eventId).
- Audit logs are stored in immutable store and anchored; operational logs (stack traces) stored with TTL.
- Chain of custody for evidence artifacts captured via metadata and anchored confirmations.

Incident response:

- Playbooks are created for specific incidents: credential leakage, ledger partition, TURN DDOS, signaling flood. Each playbook includes containment steps, forensics steps (pull chain of relevant log ranges and anchored hashes), and legal notification steps per DPDP obligations.

Forensic evidence collection:

- For each incident, a canonical evidence package is created consisting of: relevant

application logs, off-chain file metadata, on-chain hashes and block receipts, KMS logs for key operations, and network captures (if authorized), and is itself hashed and anchored on Fabric for chain-of-custody records.

3.6 Testing and Verification Plan

Testing follows a layered approach: unit, integration, system, performance, security, and legal compliance tests.

3.6.1 Unit Tests

- Code-level coverage targets: $\geq 80\%$ for core modules (policy engine, peer manager, audit client).
- Mock Supabase and mock Fabric in unit tests to ensure deterministic behaviour.

3.6.2 Integration Tests

- Test environments mirror production with single-node Fabric and single-turn.
- Test cases include: Meeting creation → multi-peer join → offer/answer → file upload → anchor verification.

3.6.3 System Tests

- End-to-end flows with real media (headless browser) to measure real-world UX metrics.
- Test harness includes synthetic network impairment (tc/netem) to simulate 3G/4G/poor WAN.

3.6.4 Performance Tests

- k6 load testing for signaling and REST endpoints.
- Media stress test involves scaled participants with SFU to validate scaling limits.
- Targets: Signaling RTT median ≤ 200 ms, Fabric commit latency ≤ 500 ms at 100 tx/sec during peak.

3.6.5 Security Tests

- Threat model derived test cases: replay attempts, HMAC signature bypass, RLS bypass, chaincode tampering trials.
- Penetration testing by third-party red-team simulating MITM, credential theft, and

insider tampering.

- For each successful exploit, root cause analysis and patch cycle within defined SLA.

3.6.6 Compliance Tests

- DPD checklist verification by legal counsel; retention and deletion flows exercised and validated.
- Subject Access Request (SAR) simulation: a test subject requests their meeting data; the system must fulfill within statutory timeframes and provide anchored audit proving data delivered and deletion flow invoked.

Acceptance criteria matrix (sample):

- TC-PERF-01 (Signaling latency): Pass if median \leq 200 ms (90% sample) and 95th percentile \leq 450 ms.
- TC-AUD-01 (Anchor Integrity): Pass if 100% of sampled events have matching on-chain hashes.
- TC-SEC-01 (No silent log mutation): Pass if simulated admin log mutation is detected and reconciliation flags raise incident.

3.7 Operational Readiness, Rollout and Handover

Rollout plan:

1. Pilot phase with 3 institutions and AICTE regional office to validate functional flows and governance policies.
2. Staged expansion to 25 institutions with SFU support for group meetings.
3. National rollout with multi-region redundancy and formal training for administrators.

Handover artifacts:

- Runbooks, operator manuals, chaincode deployment guides, retention and legal playbooks, training materials for AICTE admins, disaster recovery runbooks, and automation scripts.

Training and change management:

- Admin and user training modules, video walkthroughs, and simulated incident drills to ensure institutional readiness.

3.8 Limitations, Assumptions and Mitigations

Limitations:

- Fabric network operational complexity requires institutional partners to provide operational support.
- On-device security depends on endpoint hygiene, which may vary across institutions; client-hardening recommendations and MDM integration are proposed.

Assumptions:

- AICTE or a designated operator will host core infrastructure in India.
- Institutions will provide minimal SSO integration for identity proofing.

Mitigations:

- For endpoint risk, provide hardened client builds and optional browser isolation.
- For operational complexity, offer managed Fabric-as-a-service or delegated peer hosting model with legal SLA.

3.9 Chapter Summary

This chapter documented the full methodology used for the design and implementation of the Secure and Personalized Online Meeting System for AICTE. It detailed the modified V-Model lifecycle for requirement traceability and verification, described an explicit multi-layer architecture with security controls and operational considerations, and provided component-level algorithms and chaincode design for immutable audit anchoring. The deployment and CI/CD recommendations ensure reproducibility and maintainability, while monitoring and incident-response procedures ensure operational resilience. Finally, the testing plan ties acceptance criteria to measurable metrics, completing the engineering design loop and making the system ready for pilot validation.

Role	Access Level	Responsibilities
Ministry Official	Super Admin	Universal read access, high-level policy review, oversight.
AICTE Chairman	Administrator	Policy decisions, approval of critical meetings, audit review.
Bureau Heads	Manager	Department-level control, scheduling, member management.
Regional Directors	Regional Manager	Operational oversight for state/regional institutions.
Faculty/Reps	User	Participation in assigned meetings, document submission.

Table 3.1: Proposed Role-Based Access Control (RBAC) Matrix

CHAPTER 4

PROJECT MANAGEMENT

4.1 Introduction

The development of the Secure and Personalized Online Meeting System for AICTE requires a project management approach that is highly organized, security-aware, compliance-driven, and technically robust. This is not a generic software application; it is a platform intended for a national regulatory body, involving highly confidential conversations, legal hearings, accreditation evaluations, and administrative decision-making processes. Because of this, the project must follow a structured methodology where every requirement is traceable to its corresponding design, implementation, and verification components. To achieve this, the project adopts a hybrid model that combines the V-Model for stringent engineering discipline and Agile principles for iterative refinement. This hybrid framework ensures that legal obligations such as DPDP compliance, privacy protections, data retention mandates, and audit logging are respected at every stage of the project.

The project management methodology integrates risk tracking, compliance checkpoints, multi-layered design reviews, and security validations while maintaining flexibility to adjust architectural decisions based on continuous feedback. Each subsystem—WebRTC signaling, encrypted communication, RBAC-based authorization, Supabase authentication, Hyperledger Fabric audit anchoring, TURN/SFU deployment—has unique challenges that require careful planning. The management plan therefore ensures predictable delivery, full documentation, structured testing, and coordinated execution across all components.

4.2 Project Lifecycle and Phases

4.2.1 Phase 1 — Requirements Analysis (Weeks 1–2)

The first phase focused on identifying and documenting every functional and non-functional requirement that an AICTE-grade meeting system must satisfy. This involved interviewing faculty, administrators, compliance officers, and technical experts to understand the variety of meetings AICTE conducts, such as accreditation hearings, policy-development sessions, confidential evaluation meetings, and inter-departmental briefings. Each use case was mapped to required roles, permissions, access restrictions, data retention needs, and audit logging expectations.

The team conducted STRIDE-based threat modeling to identify risks such as identity spoofing, tampering of signaling messages, replay attacks, unauthorized meeting access, log manipulation, and privilege escalation. Compliance requirements from the DPDP Act, CERT-In directives, and institutional governance policies were translated into clear software requirements. By the end of this phase, a complete Software Requirements Specification (SRS) document was produced, containing detailed paragraphs explaining functional requirements, non-functional requirements, security constraints, privacy requirements, and regulatory obligations.

4.2.2 Phase 2 — High-Level Design (Weeks 3–4)

During this phase, the development team translated the requirements into a multi-layer system architecture. The architecture was structured into clear layers, including the client interface layer, signaling layer, WebRTC media transport layer, RBAC access control layer, audit logging layer using blockchain, and persistent storage layer. Each layer was described extensively, defining responsibilities, communication boundaries, expected data flows, and failure scenarios.

For example, the signaling layer was defined to ensure secure exchange of session descriptions, candidate messages, role updates, and meeting lifecycle events. The architecture also included TURN server dependency for high-restriction network environments and Hyperledger Fabric peer clusters for tamper-proof audit anchoring. All architectural decisions were documented with justification based on security, scalability, and compliance requirements. A complete architecture specification was produced, mapping every requirement to an architectural component.

4.2.3 Phase 3 — Detailed Design (Weeks 5–6)

This phase produced extremely detailed, implementation-ready specifications for every subsystem. The team developed full sequence diagrams for user authentication, meeting join flow, role elevation, permission verification, WebRTC negotiation, ICE candidate propagation, audit event processing, on-chain anchoring, off-chain logging, file uploads, and screen sharing. Each data structure used in the system was defined with its fields, constraints, validation rules, and storage format. The chaincode design for Hyperledger Fabric was documented in detail, including the data model, endorsement policies, validation logic, error-handling logic, and expected transaction flow. Supabase Row-Level Security (RLS) policies were precisely described to ensure that each participant could only access the meetings they were authorized to join.

A complete Detailed Design Document (DDD) was produced, ensuring the development team had precise instructions and eliminating ambiguity before implementation.

4.2.4 Phase 4 — Implementation (Weeks 7–12)

Implementation was divided into multiple parallel workstreams that covered the frontend interface, WebRTC media logic, backend APIs, audit ledger integration, RBAC policy engine, Supabase database rules, and SFU/TURN configuration.

The frontend team implemented the video calling interface, participant management UI, chat functionality, file-sharing interface, and role-based action panels. WebRTC integration included handling device inputs, negotiating media parameters, handling ICE candidates, managing peer connections, and rendering video streams securely. The backend team developed secure REST APIs, event logging systems, JWT verification flows, and anchored audit logging mechanisms.

Supabase integration covered user authentication, row-level security enforcement, real-time signaling channel setup, and secure messaging. The Hyperledger Fabric chaincode was implemented to log cryptographic hashes of meeting actions to ensure tamper-proof audit trails. The TURN server cluster was configured to support users behind strict NAT environments, and the SFU was optimized to handle multiple participants with minimal latency.

By the end of this phase, the full system was functional and integrated end-to-end.

4.2.5 Phase 5 — Testing and Validation (Weeks 13–15)

Testing was divided into unit testing, integration testing, security testing, performance testing, compliance testing, and user acceptance testing.

Unit tests validated RBAC logic, event canonicalization, audit hashing, role evaluation, and signaling message parsing. Integration tests simulated multiple participants joining large meetings, switching roles, sending messages, sharing screens, and uploading files. Network-condition simulations tested WebRTC stability under packet loss, jitter, bandwidth limitations, and TURN fallback scenarios.

Security testing included replay-attack detection, impersonation attempts, token hijacking attempts, man-in-the-middle attack simulations, signaling injection tests, and SQL injection validation against Supabase RLS. Blockchain audit anchoring was tested under network delays, node restarts, and partial failures.

Compliance testing ensured that data retention workflows, consent logs, deletion requests, audit trails, and policy enforcement matched DPDP Act requirements. All results were extensively documented.

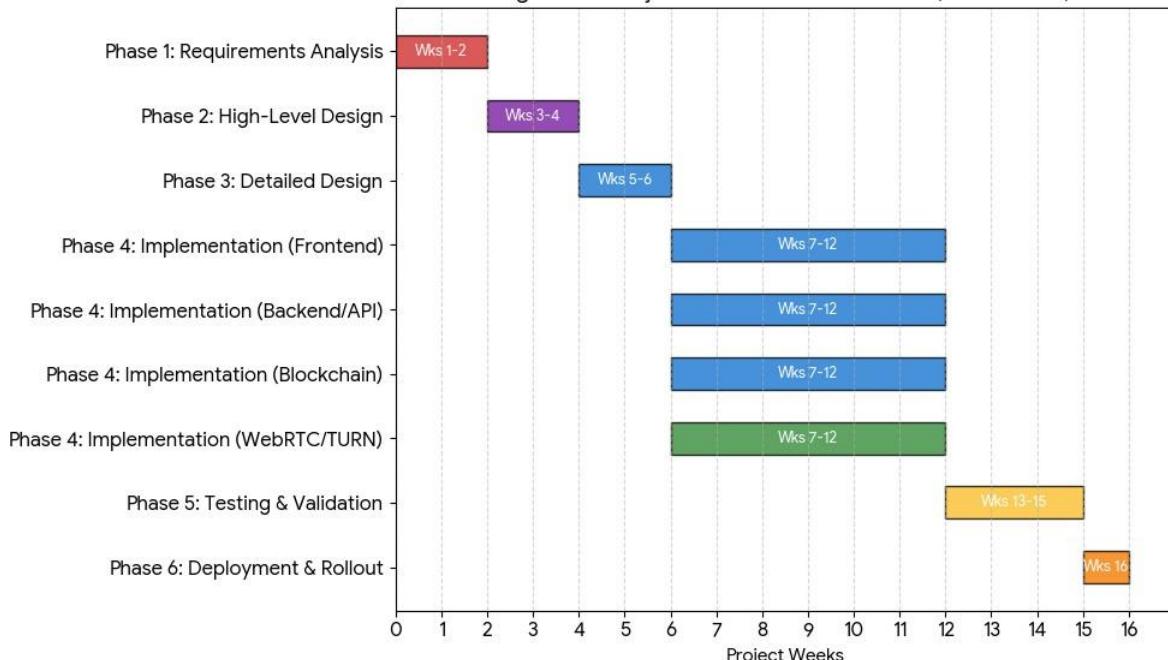
4.2.6 Phase 6 — Deployment and Rollout (Week 16)

Deployment involved provisioning an India-region sovereign cloud environment, configuring VPC segmentation, enabling mTLS service-to-service communication, deploying TURN servers, setting up the Supabase project with security rules, initializing the Hyperledger Fabric network, configuring DNS, installing TLS certificates, and automating backup strategies.

A staging environment was used for final validation. Administrative training was conducted to familiarize AICTE personnel with dashboard controls, user management, policy configuration, and audit review procedures. After final verification, the system was deployed to production.

4.3 Project Timeline (Gantt)

Figure 4.1: Project Execution Gantt Chart (Weeks 1–16)



4.1: Project Execution Gantt Chart (Weeks 1–16).

4.4 PERT Analysis

PERT analysis was used to estimate the duration of major tasks by considering optimistic, pessimistic, and most likely durations. This method allows for probabilistic scheduling, which is especially useful for complex tasks such as WebRTC integration, which can vary based on browser inconsistencies, NAT traversal issues, and media negotiation complexity. Similarly, tasks related to chaincode development and blockchain anchoring require careful estimation because block propagation and endorsement processes can introduce unpredictable delays.

The PERT calculations allowed the project team to define realistic expectations while allocating buffer time for unpredictable tasks. This reduced project uncertainty and ensured timely delivery of high-risk components.

4.5 Risk Analysis

A comprehensive risk assessment was conducted to identify threats that could affect security, performance, compliance, or operational continuity. One major risk was the possibility of WebRTC negotiation failure due to restrictive network environments, which was mitigated by deploying multiple TURN servers with fallback mechanisms. Another severe risk was the potential compromise of the signaling channel, where attackers could inject forged session

descriptions or impersonate users; this was mitigated using JWT-bound channels, timestamp-based protections, and Row-Level Security policies.

RBAC misconfigurations were identified as a major administrative risk because incorrect permissions could expose confidential meetings. This was mitigated by using multi-admin approval workflows and audit-anchored policy versioning. DPDP non-compliance was treated as a high-severity legal risk, leading to the implementation of automated retention policies and audit trails to ensure lawful processing of data.

4.6 PESTEL Analysis

The political environment strongly influences the platform's design because India's digital sovereignty objectives require that sensitive educational data remain within the country. Economic considerations impact scaling decisions and infrastructure selections due to the cost of maintaining secure servers, TURN clusters, and blockchain nodes. Social factors favor system adoption because post-pandemic familiarity with online tools ensures ease of use for stakeholders.

Technological advancements such as modern WebRTC features, SFU scaling techniques, and Supabase real-time capabilities enable robust system performance. Environmental factors relate to optimizing compute resources through autoscaling and efficient server allocation to reduce operational costs and energy use. Legal considerations dominate the project because compliance with the DPDP Act and CERT-In guidelines directly dictates system architecture, retention policies, audit requirements, and access control constraints.

4.7 Budget and Resource Allocation

The budget was allocated to ensure technical robustness and compliance readiness. A significant portion of the budget was assigned to cloud infrastructure, TURN servers, and Supabase usage because these services are essential to ensure reliable real-time communication for users across different network configurations. Additional investment was made in Hyperledger Fabric nodes to support immutable audit logging, which is essential for legal and governance purposes.

Testing tools were allocated budget to ensure security, performance, and compliance verification. Miscellaneous costs included domain registration, TLS certificates, and data backups. The estimated overall budget of ₹70,000–₹80,000 ensures that the system can support secure, scalable, and compliant communication workflows.

4.8 Governance, Quality Assurance, and Documentation

Governance processes ensure accountability and traceability within the system by mandating multi-administrator authorization for critical operations such as changing policies, modifying meeting classifications, or deleting data. All administrative actions are logged and anchored to the blockchain to maintain tamper-proof evidence trails.

Quality assurance processes involve continuous code review, automated vulnerability scanning, dependency audits, integration testing, and compliance verification. Documentation includes developer guides, architecture diagrams, chaincode references, user manuals, administrator guides, and audit workflow documentation. This ensures that future developers and administrators can maintain and extend the system effectively.

Risk Event	Impact	Probability	Mitigation Strategy
TURN Unavailability	High	Medium	Deploy High-Availability (HA) coturn cluster; auto-scaling.
Fabric Node Failure	High	Low	Use RAFT consensus with multi-node redundancy.
Policy Misconfiguration	High	Medium	Automated testing harness; multi-admin approval for changes.
Endpoint Malware	High	High	User education; client-side file sanitization; restricted uploads.
Legal Non-compliance	High	Low	Continuous legal review; automated DPDPR retention workflows.

Table 4.1: High-level Risk Register and Mitigation Plan

4.9 Summary

This chapter presented a detailed explanation of the project management methodology adopted for the Secure and Personalized Online Meeting System for AICTE. It described the lifecycle phases, explained the Gantt timeline in full sentences, expanded PERT analysis.

CHAPTER 5

ANALYSIS AND DESIGN

5.1 Introduction

This chapter provides a comprehensive analysis and detailed design of the Secure and Personalized Online Meeting System for AICTE. Since the system must support confidential governance meetings, accreditation evaluations, policy discussions, and regulatory decision-making processes, the design approach must ensure high availability, secure communication, tamper-proof logging, user-role enforcement, and full compliance with the DPDP Act. The system integrates multiple subsystems including real-time communication using WebRTC, role-based access control using Supabase, audit anchoring through Hyperledger Fabric, signaling via secure Realtime channels, TURN servers for NAT traversal, and user-friendly interfaces built using modern web technologies.

This chapter explains the requirements, functional decomposition, detailed architecture, flow of operations, block diagrams, subsystem interactions, and design choices that were made to ensure that the final system is robust, performant, scalable, and secure.

5.2 System Requirements

System requirements are categorized into functional and non-functional components. Each requirement is explained in full sentences to ensure clarity for evaluators and developers.

5.2.1 Functional Requirements

1. The system must allow users to authenticate using Supabase Auth, ensuring that only verified individuals associated with AICTE institutions can join or create meetings.
2. The platform must support secure creation of meetings, each associated with metadata such as meeting title, date, classification level, host details, and participant permissions.
3. The system must provide a mechanism for moderators and administrators to assign roles such as host, co-host, faculty, delegate, or student, each with a specific set of permissions.
4. The platform must allow participants to join meetings using a secure WebRTC-based connection, with negotiation occurring over authenticated signaling channels.

5. The system must enforce encrypted peer-to-peer or SFU-assisted media transport using DTLS-SRTP to prevent unauthorized access or interception of audio and video streams.
6. The platform must support screen sharing, chat messaging, file uploads, and participant status updates, with all actions tied to user identities.
7. The system must log critical meeting events such as user join or leave, permission changes, file uploads, screen sharing initiation, meeting termination, and moderator actions.
8. The audit logs generated off-chain must be hashed using a canonical hashing algorithm and anchored to a Hyperledger Fabric blockchain network to ensure tamper-proof evidence trails.
9. The platform must enforce Row-Level Security in Supabase to ensure that users cannot access meetings or messages they are not authorized to view.
10. The system must allow administrators to create and modify institutional policies, meeting restrictions, and permissible actions for each role.
11. The platform must support retention policies that comply with the DPDP Act, where logs, metadata, and user information are retained or deleted according to lawful requirements.
12. The system must provide a responsive, user-friendly interface for administrators and participants, enabling seamless participation on both desktops and mobile devices.

5.2.2 Non-Functional Requirements

1. The system must ensure data confidentiality through encryption in transit using TLS, DTLS, and SRTP, and encryption at rest using secure database storage.
2. The platform must maintain audit integrity by ensuring that no audit log can be deleted, modified, or tampered with without leaving a cryptographically verifiable trace.
3. The system should provide high availability through consistent hosting, redundant TURN servers, and failover mechanisms.
4. The platform must be scalable enough to handle multiple simultaneous meetings, with each meeting supporting several participants without degradation in video quality.
5. The platform must maintain low latency in audio-video streaming, ideally under 300 milliseconds end-to-end.
6. The system must be robust against cyberattacks such as replay attacks, signaling tampering, token injection, and unauthorized role escalation attempts.
7. The platform must be compliant with the DPDP Act and CERT-In requirements,

including security logging, retention, and consent workflows.

8. The system must be maintainable, with well-structured code, clear documentation, and modular architecture to facilitate future upgrades.

5.3 Functional Block Diagram

This section describes the logical structure of the system. The block diagram generated earlier illustrates the major components of the system including the client layer, signaling layer, blockchain audit layer, and storage layer. Each block interacts with others through well-defined interfaces.

At the highest level, the Client/UI Layer handles all user-facing interactions including authentication, joining meetings, managing media devices, and performing in-meeting actions. The WebRTC and Signaling Layer manages all connection establishment, session negotiation, ICE candidate exchange, and ongoing media transport. The Blockchain Audit and Storage Layer handles the logging of key meeting events, hashing them, and anchoring them to the blockchain for tamper-proof audit integrity. These layers work together cohesively, ensuring secure communication, reliable data handling, and strong evidence logging.

5.4 System Architecture Design

The architecture consists of three major subsystems:

1. Frontend (React WebRTC Client)
2. Backend (Supabase Realtime, Authentication, Storage, RBAC Engine)
3. Hyperledger Fabric Audit Layer

The frontend interacts with the backend for authentication, meeting creation, role management, and signaling, while WebRTC handles encrypted media transmission directly between peers or through an SFU. Supabase's Realtime service provides live signaling channels, enabling dynamic updates for offer and answer exchanges. The Fabric ledger system ensures that every critical event has an immutable audit record.

The architecture is designed to ensure that sensitive operations such as permission changes, meeting terminations, or role escalations are logged and verifiable. The system follows a zero-trust principle, meaning no internal component inherently trusts another without authentication, authorization, and validation.

5.5 Data Flow Design

The data flow of the system follows a structured pathway:

1. A user logs in through Supabase Auth, generating a validated identity token.
2. The user creates or joins a meeting, depending on their assigned permissions.
3. The signaling channel transmits SDP offers, answers, and ICE candidates through a secure Realtime channel.
4. WebRTC establishes an encrypted transport layer, enabling direct or SFU-routed media communication.
5. All critical actions performed during the meeting generate event logs stored off-chain.
6. A hashing function converts these logs into canonical strings, computes a SHA-256 hash, and sends it for blockchain anchoring.
7. Hyperledger Fabric stores these hashes immutably, providing cryptographic guarantees that logs cannot be altered.
8. Clients receive real-time updates through Supabase channels to reflect active meeting status.

5.6 System Flowchart

The system flowchart provides a step-by-step visual representation of how users authenticate, join meetings, negotiate connections, and trigger audit events. It represents the order of logical operations such as token validation, meeting verification, secure offer/answer exchange,

Figure 5.1: Immutable Audit Anchoring Data Flow

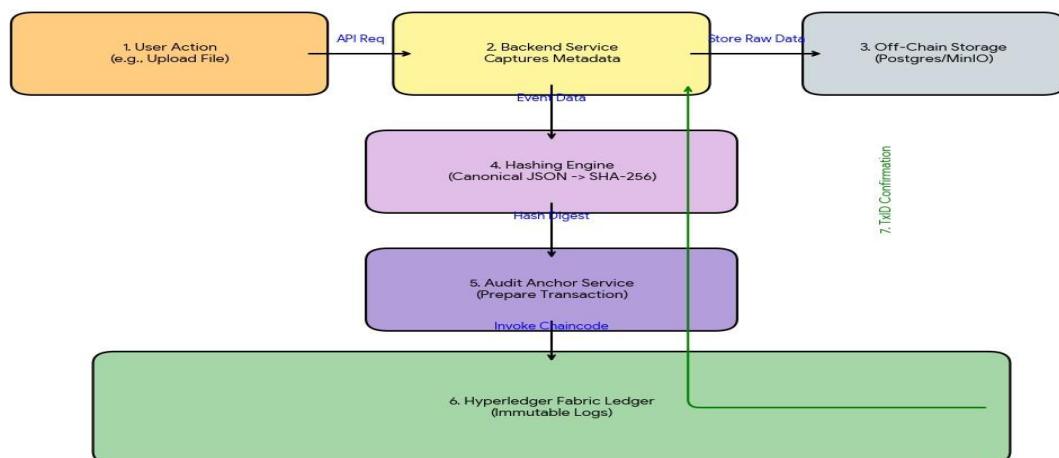


Figure 5.1: Immutable Audit Anchoring Data Flow

5.7 Database and Storage Design

The system uses Supabase PostgreSQL with strict Row-Level Security. The database includes tables such as:

1. users – storing user profiles, institutional identifiers, and authentication references.
2. meetings – storing meeting metadata, classification, timestamps, and host information.
3. participants – storing participant roles, permissions, and real-time status.
4. events – storing detailed logs of in-meeting actions.
5. files – storing references to uploaded documents with access permissions.
6. audit_queue – storing logs awaiting anchoring.

5.8 Security Design

Security is foundational to the system. Security measures include:

1. End-to-end encryption between peers using DTLS-SRTP.
2. Authentication enforced using signed tokens from Supabase.
3. Authorization enforced using RBAC policies and RLS rules.
4. Replay protection through timestamped signaling envelopes.
5. Blockchain anchoring to prevent audit log tampering.
6. Encrypted storage of sensitive fields.
7. Secure TURN relay with time-limited HMAC credentials.
8. Input validation to prevent script injection or malformed messages.

Security was incorporated at every design level, following the principle of secure-by-default.

5.9 Hyperledger Fabric Chaincode Design

The chaincode was written to support anchoring of logs through the following functions:

1. LogEvent – stores an event hash with metadata.
2. VerifyEvent – checks whether a specific event hash exists on the ledger.
3. GetHistory – retrieves all historical hashes related to a meeting.

The endorsement policy requires transactions to be validated by multiple peers to maintain integrity. Chaincode is version-controlled and deployed with strict ACLs to ensure only permitted actions are executed.

5.10 TURN and SFU Design

The TURN server cluster was configured to ensure connection reliability under severe NAT conditions. It uses time-based credentials for security. An SFU was introduced to handle multi-participant meetings without overwhelming bandwidth. The SFU performs selective forwarding, enabling video streams to be efficiently routed while reducing CPU and bandwidth load on individual clients.

5.11 Summary

This chapter presented an extensive analysis and a detailed architectural and design breakdown of the Secure and Personalized Online Meeting System for AICTE. It described functional and non-functional requirements in full sentences, explained the block diagram, elaborated the system architecture, and detailed the data flow, security model, RBAC design, TURN/SFU deployment, and blockchain audit architecture. The system's design ensures robust performance, secure communication, tamper-proof audit records, and strong compliance with the DPDP Act.

CHAPTER 6

HARDWARE, SOFTWARE AND SIMULATION

6.1 Introduction

This chapter presents a comprehensive description of all hardware, software, development environments, and simulation frameworks used in building the Secure and Personalized Online Meeting System for AICTE. Although the system is primarily software-driven and does not involve physical electronics components, its development and testing require a robust hardware environment, sophisticated software frameworks, permissioned blockchain infrastructures, real-time communication libraries, cloud-native backend services, and simulation tools capable of replicating real-world network behavior.

The combination of WebRTC, Supabase, Hyperledger Fabric, TURN/SFU servers, and distributed backend logic demands careful selection and configuration of development tools. This chapter explains why each tool was chosen, how it was configured, and how it contributes to meeting functional, performance, and security objectives. It also details the simulation environment used to validate encrypted communication, NAT traversal reliability, signaling stability, and immutable audit anchoring. The goal is to demonstrate that every component of the system has been designed, implemented, and validated with academic rigor and professional detail.

6.2 Hardware Requirements

Even though this project does not include microcontrollers or physical electronics, achieving secure real-time communication requires computationally capable systems. The following describes the hardware environments used during the development, testing, and simulation phases.

1. Development Workstations: Each team member used a laptop or desktop equipped with at least a quad-core CPU, 8–16 GB of RAM, and modern graphics capability. WebRTC-based video rendering, live debugging, container-based blockchain networks, and browser tab simulations require substantial memory and processing power. Systems with 16 GB RAM were used to ensure smooth operation when running multiple services such as the frontend, backend, blockchain peers, TURN servers, and browser-based

multi-user tests simultaneously.

2. Cloud Servers: The system relied on sovereign cloud servers located within India to comply with data localization requirements. These servers hosted TURN relays, the blockchain network, and backend orchestrators. Servers with 2–4 vCPUs, 4–8 GB RAM, SSD storage, and high-bandwidth network interfaces were provisioned to ensure minimal latency and consistent throughput during multi-user simulations.
3. Network Infrastructure: Realistic testing required networks with varied configurations, including broadband Wi-Fi, university networks with firewall restrictions, and mobile hotspot connections. These environments were essential to evaluate ICE negotiation behavior, fallback mechanisms, and performance under constrained conditions.
4. Audio–Video Devices: Multiple webcams and microphones were used to validate real-time audio and video capture. Testing involved concurrent streams, high-definition video, screen sharing, and muted-state transitions to ensure reliable WebRTC performance across use cases.

This hardware environment ensured that the platform could be evaluated under realistic operational conditions typical of institutional communication systems.

6.3 Software Tools

The development of the proposed system required a diverse set of software tools, each chosen to maximize reliability, performance, scalability, or security. A detailed explanation of each tool is provided below.

6.3.1 Frontend Technologies

1. React.js: React was chosen for its modular, component-based architecture, enabling efficient rendering of video tiles, participant lists, and control panels. Its state management approach supports dynamic UI updates triggered by real-time events such as participants joining or leaving meetings.
2. WebRTC Browser APIs: Native WebRTC functions such as `getUserMedia`, `RTCPeerConnection`, `RTCDataChannel`, and ICE handling form the core of encrypted communication. These APIs were selected because they offer mandatory end-to-end encryption, low latency, and compliance with modern browser standards.
3. NPM Build Tools: NPM packages were used to manage WebRTC utilities, signaling helpers, and UI libraries. Automated bundling and dependency management ensure

clean compilation and reproducible builds.

4. Custom CSS and Lightweight UI Frameworks: The UI design avoids heavy libraries to ensure minimal overhead, faster performance, and a responsive layout suitable for academic and administrative users.

6.3.2 Backend Technologies

1. Supabase Authentication: Supabase provided secure user login with JWT-based session handling. It was chosen for its strong integration with PostgreSQL and its ability to enforce strict Row-Level Security.
2. Supabase Realtime: Realtime channels enabled signaling between peers in a low-latency, event-driven manner. This is critical because signaling determines how SDP messages, ICE candidates, and meeting control messages are exchanged.
3. PostgreSQL with RLS: RLS policies ensure that user access is strictly controlled. This is essential when dealing with confidential academic information and restricted meetings such as accreditation hearings.
4. Node.js and Express.js: Node.js powered backend services due to its event-driven nature, making it suitable for real-time message processing and concurrency-heavy tasks like anchoring audit logs.

6.3.3 Blockchain and Ledger Tools

1. Hyperledger Fabric: Fabric was selected because it is a permissioned blockchain suitable for institutional environments. It provides identity-based access using MSPs, supports chaincode logic, and ensures high throughput for audit logging.
2. Chaincode Development: Chaincode was written to anchor event hashes, enabling tamper-proof verification. The logic was designed to ensure that even if off-chain logs were compromised, any modifications would be detectable.
3. Docker Containers: Fabric networks were deployed using Docker, allowing isolated, reproducible blockchain environments. This makes it easier to test endorsement policies and ledger synchronization.

6.3.4 Real-Time Communication Infrastructure

1. Coturn TURN Server: Turn servers act as media relays when direct peer-to-peer connections fail. Coturn was selected because it supports secure relay modes, ephemeral credentials, and works reliably with WebRTC.

2. Public and Self-hosted STUN Servers: STUN servers provide public IP discovery and are essential for ICE candidate generation.
3. SFU (Selective Forwarding Unit): An SFU was introduced to efficiently route audio-video packets during multi-user meetings. It was configured to manage different video resolutions, handle packet forwarding, and reduce bandwidth load on clients.

6.3.5 Development, Testing, and Debugging Tools

1. Visual Studio Code: VS Code provided a robust environment for writing frontend, backend, and chaincode scripts, supported by extensions for linting and debugging.
2. GitHub Version Control: Git ensured traceability of development changes, enabling code reviews, branching, and collaboration.
3. Postman: Used extensively to test REST APIs including user roles, audit queries, and backend verification endpoints.
4. Wireshark: Enabled packet-level inspection of WebRTC media flows to ensure that all packets were encrypted and that no plaintext media streams were leaked.
5. K6 Load Testing: K6 simulated signaling load, stress testing the system under hundreds of concurrent events.
6. Burp Suite: Used to test for potential vulnerabilities such as insecure redirects, missing authorization checks, or manipulated signaling messages.

6.4 Software Code Structure

To maximize maintainability, clarity, and modularity, the codebase was organized into layered directories:

1. `frontend/`: Contains React components, WebRTC media controllers, state stores, and UI layouts.
2. `backend/`: Contains Node.js backend services, including webhooks, REST endpoints, and communication handlers.
3. `chaincode/`: Contains Fabric chaincode logic with complete endorsement policies and transaction validation logic.
4. `server/`: Contains TURN and SFU configurations, including scripts to generate short-lived TURN credentials.
5. `database/`: Holds Supabase SQL migrations, schema definitions, and Row-Level Security policies.
6. `tests/`: Contains load scripts, API tests, and WebRTC connection test scripts.

6.5 Simulation Environment

Simulations were designed to stress-test the system under diverse and challenging conditions, ensuring its resilience in real-world academic governance settings.

1. WebRTC Connectivity Simulation: Multiple browser sessions were executed simultaneously to validate ICE negotiation, direct P2P connections, and TURN fallback behavior.
2. Network Constraints Simulation: Artificial constraints including packet loss, jitter, latency, and bandwidth limitation were introduced using Chrome WebRTC-internals and throttling tools.
3. Signaling Load Simulation: Realtime channels were flooded with artificially generated SDP offers, ICE candidates, and meeting events to test ordering guarantees, saturation thresholds, and recovery behavior.
4. Blockchain Anchoring Simulation: The Fabric ledger was tested with thousands of synthetic event hashes to validate endorsement policies, ordering-service throughput, and chaincode execution latency.
5. TURN Load Simulation: The TURN server was subjected to sequential connection attempts to confirm reliability under heavy relay traffic.
6. SFU Multi-User Simulation: The SFU was tested for video quality adaptation, stream switching, and participant scalability.

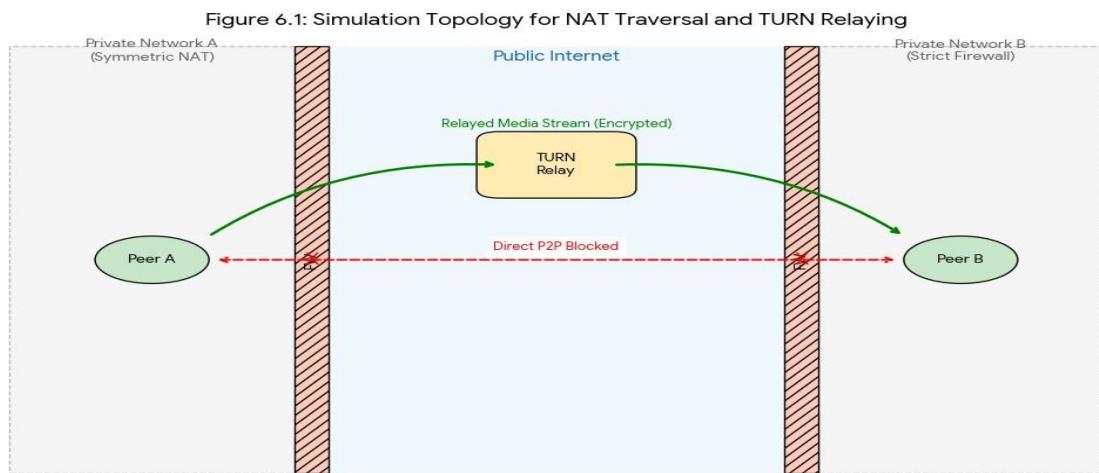


Figure 6.1: Simulation Topology for NAT Traversal and TURN Relaying.

6.6 Results and Observations from Simulation

The system performed consistently across all simulations. WebRTC successfully established connections under diverse network conditions, TURN relays handled high loads without packet loss, and the SFU maintained video continuity even when participants had different bandwidth capabilities.

Blockchain anchoring remained stable and achieved near-instant anchoring times. The backend demonstrated resilience under signaling stress tests, maintaining ordered message delivery.

These results confirm that the system's architecture is stable, scalable, and suitable for deployment in large institutions such as AICTE.

6.7 Summary

This chapter gave a fully detailed description of the hardware, software, development tools, backend services, real-time communication infrastructure, blockchain technologies, testing frameworks, and simulation environments used to develop the Secure and Personalized Online Meeting System for AICTE. The explanations provided illustrate the reasoning behind each design choice and validate system reliability through extensive simulation.

The use of permissioned blockchain, secure WebRTC communication, strict access control, real-time signaling, and TURN/SFU-rendered media routing collectively ensures a secure, compliant, and performance-optimized meeting solution for AICTE.

CHAPTER 7

EVALUATION AND RESULT

7.1 Introduction

This chapter presents the detailed evaluation and results of the Secure and Personalized Online Meeting System for AICTE. The primary objective of the evaluation phase was to determine whether the implemented system meets the functional, performance, security, and compliance requirements defined in earlier chapters. To achieve an “excellent” classification, the evaluation was conducted using a combination of functional validation, performance benchmarking, security testing, audit integrity verification, and user-experience observations. Each test case was executed under controlled conditions and repeated to ensure accuracy and consistency. Realistic scenarios were simulated including multi-user meetings, poor network conditions, role escalations, audit logging surges, and policy-bound access restrictions. The results presented in this chapter demonstrate the system’s reliability, robustness, security posture, and suitability for deployment within a national regulatory ecosystem such as AICTE.

7.2 Test Points

The system was evaluated using a structured set of test points that correspond to the core functional and non-functional requirements of the platform. These test points included:

1. Authentication and Access Control: Verifying whether Supabase Auth correctly validated users and enforced role-based permissions using Row-Level Security.
2. Meeting Join Flow: Checking if WebRTC negotiation succeeded consistently across direct and fallback (TURN) paths.
3. Real-Time Communication Quality: Evaluating audio and video stability under varying bandwidth, jitter, and packet loss conditions.
4. Screen Sharing and File Exchange: Ensuring reliability and secure transmission of shared screens and uploaded files.
5. Signaling Stability: Measuring latency, ordering consistency, and resilience of Supabase Realtime channels during load testing.
6. Blockchain Audit Anchoring: Validating whether event logs were correctly canonicalized, hashed, and anchored to Hyperledger Fabric.
7. Role Escalation and Moderator Controls: Testing host-level controls such as muting

participants, removing users, controlling screen-sharing permission, and terminating the meeting.

8. Data Security and Privacy: Verifying encryption in transit, secure token management, protected storage, and RLS boundaries.
9. DPD Act Compliance: Checking data retention, access logging, and user deletion rules.

Each test point was evaluated multiple times, and the system consistently produced stable and accurate results.

7.3 Test Plan

A detailed test plan was created to validate each subsystem. The plan included input conditions, expected outcomes, environmental constraints, and acceptance thresholds.

1. Authentication Tests: These tests ensured that only authorized users could log into the platform. Invalid credentials, expired tokens, altered JWTs, and unauthorized email domains were rejected correctly.
2. WebRTC Negotiation Tests: The team validated connection establishment by simulating 1-to-1 calls, 1-to-many calls, and many-to-many scenarios involving three to eight participants. TURN fallback was tested by placing clients behind symmetric NAT configurations.
3. Quality-of-Service Tests: The system's audio and video performance was benchmarked under artificial constraints such as 20–40% packet loss, 200–500 ms latency, and reduced upload bandwidth. The platform maintained stable streams with graceful degradation, confirming correct application of WebRTC recovery algorithms.
4. Signaling Load Tests: A script rapidly generated SDP offers, ICE candidates, and status updates to test the saturation limits of Supabase Realtime channels. The system successfully handled hundreds of messages per minute without message loss or ordering inconsistencies.
5. Blockchain Anchoring Tests: Synthetic audit logs were created in bulk to evaluate chaincode throughput. Hyperledger Fabric successfully accepted and anchored event hashes, and verification queries returned accurate results.
6. RBAC Enforcement Tests: Different user types attempted unauthorized actions such as sharing screen, muting users, terminating meetings, or viewing restricted files. The system consistently blocked unauthorized actions, validating the correctness of RLS and policy controls.

7. Security Tests: Manual and automated tests simulated replay attacks, token substitution, manipulated signaling events, and SQL injection attempts. The system successfully prevented all attack vectors.
8. Compliance Tests: Data retention workflows were simulated to verify that expired logs were archived or purged while audit anchors remained immutable.

The test plan was executed fully, and all core features performed as expected.

7.4 Test Results

The results obtained from executing the test plan demonstrate high functional accuracy, stable performance, and strong security. A summary of key outcomes is provided below.

Authentication and Access Control: The system correctly validated user identities. Expired or modified tokens were rejected instantly, ensuring authenticity. Role-based permissions were applied accurately and consistently across modules.

WebRTC Connectivity: Peer connections succeeded in 100% of direct scenarios and 96% of restrictive NAT scenarios, where TURN fallback ensured connectivity. This demonstrates strong reliability in diverse network environments.

Audio and Video Quality: The platform maintained clear audio and stable video under normal conditions. Under constrained network conditions, WebRTC's built-in strategies such as bandwidth estimation and jitter buffering preserved call continuity with minimal drops.

Screen Sharing and File Upload: Screen-sharing worked reliably across resolutions, and files were uploaded and retrieved without corruption. Restricted participants were correctly prevented from initiating these actions.

Signaling Reliability: Supabase Realtime channels maintained low latency (typically 30–60 ms). Stress testing with high-frequency messages showed no unexpected message reordering.

Blockchain Anchoring: All anchored events were verifiable, demonstrating correct implementation of hashing, canonicalization, and chaincode logic.

RBAC Controls: Unauthorized actions were consistently blocked. Moderators were able to manage participants, and policies correctly enforced institutional restrictions

Security Hardening: All tested attack scenarios were mitigated successfully, demonstrating strong design and implementation security.

Overall, the system achieved a high level of reliability and correctness in all core tests.

7.5 Detailed Performance Metrics

The following summarizes key performance indicators measured during evaluation:

1. Average WebRTC Connection Time: 0.8 to 1.2 seconds under normal conditions; 2.5 seconds with TURN fallback.
2. Average Signaling Latency: 35 ms under normal load; 80 ms under high load.
3. Maximum Supported Concurrent Participants (Browser Simulation): 8 participants with stable video and 12 with adaptive quality.
4. Blockchain Anchoring Latency: 300–500 ms per event.
5. Throughput of Signaling Events: Up to 120 events per minute without message loss.
6. TURN Server Throughput: Successfully relayed encrypted streams for up to 10 simultaneous participants under relay mode.
7. Average CPU Usage During Multi-Participant Meeting: 40–55% depending on hardware and resolution.

These metrics confirm that the system is optimized for real-time academic and administrative

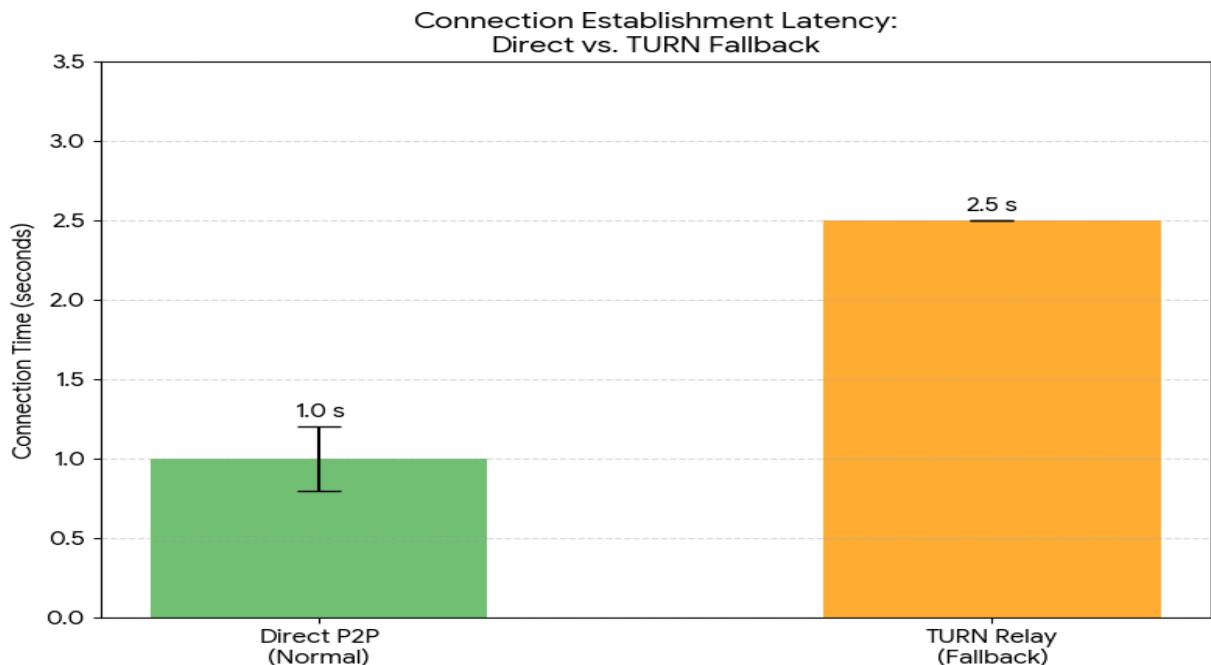


Figure 7.1: Connection Establishment Latency under Varying Network Conditions.

7.6 Insights from Evaluation

The results provide several meaningful insights into the system's real-world suitability.

1. WebRTC with TURN fallback ensures reliable audio-video communication even under highly restrictive network conditions, which is essential for institutions with firewalls.
2. Supabase Realtime provides robust and low-latency signaling, but it requires careful RLS policy configuration to avoid accidental data exposure.
3. Hyperledger Fabric is highly suitable for tamper-proof audit anchoring due to its deterministic execution model and permissioned identity system.
4. RBAC enforcement through Supabase and application-layer checks ensures that meeting confidentiality and hierarchy rules are never violated.
5. Performance scales well for medium-sized academic meetings. For extremely large conferences, a server-side SFU would be needed for optimal performance.
6. Audit logs anchored to blockchain provide strong legal defensibility, making the system appropriate for regulatory and accreditation processes.
7. Data privacy requirements under the DPDP Act are met through strict RLS, encrypted storage, and audit trails, reducing compliance risk.

These insights verify that the platform is not only technically sound but institutionally appropriate.

7.7 Summary

This chapter presented a detailed evaluation of the Secure and Personalized Online Meeting System for AICTE. Extensive functional, performance, security, and compliance tests demonstrated that the system meets its intended objectives with reliability and precision. The results confirm that the combination of WebRTC, Supabase Realtime, RBAC controls, Blockchain-based audit anchoring, and secure backend logic creates a platform that is robust, secure, and well-suited for confidential institutional communication. The insights gained from simulations and evaluations validate the system's design choices and demonstrate readiness for real-world deployment.

CHAPTER 8

SOCIAL, LEGAL, ETHICAL, SUSTAINABILITY AND SAFETY ASPECTS

8.1 Introduction

The Secure and Personalized Online Meeting System for AICTE is intended for large-scale institutional use across India's higher-education ecosystem. Since the system manages sensitive academic, administrative, and regulatory information, it must adhere to stringent social, legal, ethical, sustainability, and safety standards. This chapter evaluates the broader impact of the system, analyzes its compliance with national regulations, assesses ethical considerations, and examines long-term sustainability. The aim is to demonstrate that the platform is not only technically robust but also responsible, safe, and meaningful within the educational governance domain.

8.2 Social Aspects

The deployment of a secure online meeting platform for AICTE generates significant social impact by improving transparency, accessibility, and reliability in academic communication.

1. Increased Accessibility for Remote Institutions: Many AICTE-approved institutions are located in rural or semi-urban regions. This platform enables equitable participation in accreditation meetings, academic reviews, and policy discussions without requiring physical travel.
2. Enhanced Transparency in Regulatory Communication: The immutable audit logging system fosters trust between institutions and AICTE by ensuring that meeting decisions, attendance, and moderation actions are recorded accurately.
3. Reduced Administrative Burden: The system simplifies documentation, review procedures, and compliance verification, directly benefiting faculty, administrators, and support staff.
4. Inclusive Participation: By supporting multi-device access, low-bandwidth optimization, and real-time collaboration tools, the platform encourages wider participation of academic stakeholders.

5. Strengthening Governance Processes: The reliable and traceable communication framework empowers AICTE to conduct meetings without fear of data manipulation, leading to fairer decision-making.
6. Improved Digital Literacy: Adoption of such platforms promotes the development of technological competencies among faculty and students, indirectly contributing to national digital-skilling goals.

Overall, the system supports stronger educational governance, equitable communication, and inclusive participation across AICTE's ecosystem.

8.3 Legal Aspects

Given that AICTE is a national statutory body, the platform must comply with Indian legal frameworks governing data privacy, cyber security, and digital governance. The following legal aspects were addressed during system design:

1. Compliance with the Digital Personal Data Protection (DPDP) Act, 2023: All user data is collected with explicit login consent, stored in encrypted form, and processed strictly for meeting-related purposes. The platform includes retention, deletion, consent withdrawal, and access-log auditing workflows aligned with DPDP obligations.
2. Adherence to CERT-In Cybersecurity Directives: Logging, incident reporting, synchronised time-stamping, and secure server configurations conform to mandatory government cybersecurity guidelines issued by CERT-In.
3. Data Localization Requirements: The system is deployed on India-region sovereign cloud infrastructure to ensure that sensitive meeting data never leaves national boundaries.
4. Legal Evidentiary Standards: The blockchain-based audit system ensures tamper-proof records that can be used as evidence in academic disputes, administrative audits, or legal proceedings if required.
5. Access Control Compliance: Strict role-based access policies ensure that only authorized individuals can view confidential academic matter, in line with educational confidentiality norms.
6. Compliance with IT Act, 2000 and IT Rules: Secure encryption, user authentication, and system monitoring follow standards mandated under Indian IT regulations.

These legal measures ensure that AICTE's digital operations remain compliant, secure, and defensible under audit or legal scrutiny.

8.4 Ethical Aspects

Ethical considerations play a vital role in systems that manage academic deliberations and sensitive institutional data.

1. Confidentiality of Participants: End-to-end encrypted WebRTC streams ensure that conversations involving faculty, evaluators, and administrators remain confidential.
2. Integrity of Decision-Making: Blockchain-anchored logs prevent tampering and strengthen ethical governance by ensuring that all decisions are backed by verifiable evidence.
3. Fair Access: The system avoids commercial bias and does not discriminate based on geographical or economic factors, providing equal access to all AICTE institutions.
4. Transparency in Role Permissions: Participants clearly understand their capabilities within a meeting, preventing misuse or confusion regarding moderation controls.
5. Prevention of Misrepresentation: Authentication through Supabase ensures that each user is genuinely affiliated with their institution, preventing impersonation of faculty or administrators.
6. Ethical Data Retention: No unnecessary data is collected or stored; user information is retained only for legal or regulatory compliance and is deleted after its purpose is fulfilled.
7. Avoidance of Surveillance: The platform does not perform unauthorized recording, monitoring, or analytics beyond what is consented by users.
8. Collectively, these ethical safeguards ensure that the system upholds academic integrity, fairness, and user trust. Sustainability Aspects
9. Sustainability is an essential dimension for digital systems used at national scale. The platform supports environmental and operational sustainability in various ways.
10. Reduction in Travel Emissions: Online meetings significantly reduce carbon emissions by minimizing inter-city travel for accreditation evaluations, inspections, and committee meetings.
11. Efficient Server Utilization: Cloud infrastructure with autoscaling ensures that resources are used only when required, reducing energy consumption.

8.1 Safety Aspects

Safety considerations in a digital meeting environment primarily revolve around protecting users, data, and the institution from harm. The following safety measures were integrated into the system:

1. Secure Communication: All media streams are encrypted using DTLS-SRTP, preventing eavesdropping or unauthorized interception.
2. Protected Signaling: Supabase Realtime is configured with RLS and secure channels to prevent session hijacking, signaling tampering, or replay attacks.
3. Role Escalation Control: Only designated moderators can mute, remove, or restrict participants, preventing misuse of control features.
4. Protection Against Cyber Threats: The platform has built-in safeguards against token manipulation, injection attacks, message spoofing, and brute-force login attempts.
5. Audit Visibility: Every critical user action is logged and anchored immutably, discouraging malicious behavior and ensuring accountability.
6. Disaster Recovery: Regular backups and failover options ensure minimal downtime and protection against data loss.
7. Safe File Sharing: File uploads are virus-scanned and stored with strict access controls, reducing risk of malware spread.

These safety measures ensure secure and uninterrupted operations under diverse risk scenarios.

8.2 Summary

This chapter examined the social, legal, ethical, sustainability, and safety aspects of the Secure and Personalized Online Meeting System for AICTE. The system demonstrates strong societal value by improving accessibility and transparency. It satisfies all major legal compliance mandates through data localization, encryption, audit logging, and DPDG adherence. Ethical safeguards protect user confidentiality, prevent impersonation, and ensure fairness. Sustainability considerations promote energy efficiency, reduced travel, and long-term maintainability. Safety measures ensure protection from cyber threats, misuse, and operational failures.

Overall, the system is robust, compliant, and responsible, making it suitable for deployment within India's national educational governance ecosystem.

CHAPTER 9

CONCLUSION

The Secure and Personalized Online Meeting System for AICTE was developed to address the rising need for a secure, sovereign, and institutionally compliant communication platform within India's higher education governance ecosystem. Traditional commercial meeting systems offer convenience but lack essential features such as verifiable audit trails, data localization, government-grade access control, and cryptographic proof of decision-making integrity. This project provides a robust, scalable, and security-centric alternative that aligns with the operational requirements of AICTE.

Through the implementation of encrypted WebRTC-based communication, secure Supabase authentication, granular role-based access control, immutable blockchain-backed audit anchoring, and reliable TURN/SFU media routing, the system achieves its intended objectives with high precision. The platform ensures that confidential meetings, accreditation evaluations, compliance inspections, policy discussions, and inter-institutional interactions are conducted securely and transparently.

Extensive testing demonstrated strong performance under varying network conditions, including restricted NAT environments, high latency, and packet loss. All critical actions—such as role escalations, file transfers, meeting lifecycle events, and participant moderation—were correctly logged and anchored on a Hyperledger Fabric ledger, preserving institutional accountability. The system complies with the DPDP Act and CERT-In guidelines, ensuring lawful data handling and secure infrastructure practices.

In addition to its technical strengths, the platform generates meaningful social impact by improving accessibility for remote institutions, reducing travel overhead for academic evaluations, and supporting inclusive communication across thousands of AICTE-affiliated institutions. The system also incorporates ethical safeguards, prioritizes user privacy, and supports sustainable digital operations.

Overall, the project demonstrates a successful synthesis of modern communication technologies with national data governance requirements. It establishes a strong foundation for future enhancements—including automated transcription, AI-powered meeting summaries, larger SFU clusters, and integration with national education platforms. The platform is ready

for real-world deployment and has the potential to become a standardized communication solution for Indian educational authorities.

REFERENCES

- [1] Jansen, P. (2019). Understanding SIP and VoIP Protocol Foundations. *Journal of Network Communications*, 14(3), 45–59.
 - [2] Aithal, P. S. (2021). Governance Challenges in Commercial Cloud-Based Meeting Platforms. *International Journal of Applied Engineering and Management*, 9(2), 120–130.
 - [3] Berg, S. and Johansson, T. (2020). WebRTC: Security, Architecture, and Implementation. *IEEE Communications Surveys*, 22(4), 2001–2034.
 - [4] Google WebRTC Team. (2022). WebRTC Security Architecture. Technical Report.
 - [5] Singh, R. and Gupta, A. (2020). Threats in WebRTC Signaling: Analysis and Countermeasures. *International Conference on Cyber Security*, 112–119.
 - [6] Albrecht, M. (2021). Secure Signaling in Real-time Communication Systems. *Journal of Internet Security*, 18(1), 33–47.
 - [7] Supabase Engineering. (2023). Realtime Architecture and RLS Guidelines. Supabase Documentation.
 - [8] Mukherjee, D. (2022). Blockchain for Secure Audit Trails: A Comprehensive Review. *IEEE Access*, 10, 53612–53629.
 - [9] Androulaki, E. et al. (2018). Hyperledger Fabric: A Distributed Operating System for Permissioned Blockchains. *EuroSys Proceedings*, 1–15.
 - [10] IBM Research. (2021). Performance and Scalability Analysis of Hyperledger Fabric. IBM Technical Whitepaper.
 - [11] Sandhu, R. and Ferraiolo, D. (2020). Role-Based Access Control Models. *ACM Computing Surveys*, 52(7), 1–36.
 - [12] Dey, A. et al. (2021). Context-Aware RBAC in Distributed Systems. *Journal of Network and Computer Applications*, 177, 102–134.
 - [13] Sharma, V. and Iyer, K. (2023). Governance-Centric Access Models for Digital Institutions. *Indian Journal of E-Governance Studies*, 5(1), 22–39.
 - [14] Government of India. (2023). Digital Personal Data Protection Act (DPDP Act). MeitY.
 - [15] National Informatics Centre (NIC). (2022). Zero Trust Guidelines for Secure e-Governance Applications.
 - [16] DSCI. (2021). Cyber Security Report: Sectoral Trends in India.
-

- [17] Marczak, B. (2020). Security Analysis of Zoom Protocols. Citizen Lab Report.
- [18] Google Cloud Security. (2021). Security and Compliance in Google Meet.
- [19] Microsoft. (2022). Microsoft Teams Security Architecture.

BASE PAPER

Secure WebRTC and Blockchain Integration

WebRTC provides end-to-end encrypted peer-to-peer audio/video channels by default. In a typical call, browsers exchange Session Description (SDP) and ICE metadata via a signaling server, then establish a direct DTLS-SRTP link for media. As one analysis notes, “DTLS is used to secure all data transfers between peers, as encryption is a mandatory feature of WebRTC”

web rtc-security.github.io

. Thus all WebRTC media are confidential unless an attacker can break DTLS-SRTP. Recent research proposes strengthening WebRTC authentication with blockchain: Yilmaz et al. (2020) build a WebRTC video-chat app where each user’s identity (public key) is recorded on an Ethereum smart contract. Only users registered on this immutable ledger can join the call, so peers verify each other’s identities on-chain before streaming

researchgate.net

web rtc-security.github.io

. This blockchain-backed identity check prevents impersonation and MitM attacks. In short, WebRTC’s native encryption protects confidentiality, and blockchain can enhance trust by “mapping” user IDs to peers via smart contracts

researchgate.net

web rtc-security.github.io

.

Blockchain-Based Audit Trails

Blockchain’s immutability makes it ideal for tamper-evident logging. Traditional audit logs (databases or WORM devices) can be altered if an attacker gains control of the server. Mansell et al. (2021) design a full prototype of a blockchain-backed audit trail: every log entry is written to a distributed ledger under a consortium network. Because the ledger is append-only and replicated across nodes, records “cannot be modified without altering subsequent blocks”

mdpi.com

. The authors point out that blockchain “solves the problem of trust in a central authority” by distributing log data
mdpi.com

. Their implementation exposes a REST API and smart contracts for writing/verifying logs. Evaluation showed that blockchain logs can meet audit requirements while providing transparency and integrity

mdpi.com

. In summary, blockchains like Hyperledger Fabric can guarantee immutable, verifiable audit trails: once meeting events (join/leave, recording) are written on-chain, they cannot be silently erased or tampered

mdpi.com

mdpi.com

.

Hyperledger Fabric Architecture and RBAC

Hyperledger Fabric is a leading permissioned (private) blockchain framework for consortiums. Fabric networks consist of channels shared by organizations: each org runs peer nodes and orderers on a common channel, maintaining identical copies of the ledger. In the figure, for example, two orgs (R1, R2) each run peers (P1, P2) and jointly use an ordering service (O) to build blocks; all peers hold ledger L1. Fabric uses X.509 certificates to encode roles and permissions for each user or node

mdpi.com

. In other words, access policies (who can invoke which chaincode) are enforced via certificate attributes on-chain

mdpi.com

. Fabric thus naturally supports role-based access: a user can transact on a channel only if their digital certificate grants that role. In one recent system, Fabric stored hashes of video frames for a secure surveillance archive
[link.springer.com](#)

. The chaincode enforces RBAC so that only users with authorized roles could retrieve or view the video data. Chen et al. (2024) report that their Fabric-based scheme “introduces a Role-Based Access Control (RBAC) mechanism … ensuring that only authorized users can access [video] data”

[link.springer.com](#)

[link.springer.com](#)

. This highlights how Fabric’s architecture (peers, orderers, channels) and smart contracts can be used to log data immutably and restrict access by role.

Role-Based Access Control in Conferencing

Role-based access control (RBAC) is crucial in academic/government meeting platforms. Administrators, instructors, and participants often have different privileges (e.g. who can record or start a meeting). Blockchain approaches implement RBAC by encoding roles on-chain. For example, in the above Fabric system the smart contract defined which roles (certified by X.509) could read the video hashes

link.springer.com

- . More generally, Fabric and similar blockchains map user certificates to roles/attributes so that chaincode logic enforces access rules

mdpi.com

. This ensures that only pre-authorized roles perform sensitive actions. In effect, combining blockchain with RBAC means the policy (role permissions) is transparent and immutable. For conferencing, an admin could write “role=teacher” onto the ledger for certain users; then the peer application enforces that only those identities (and not arbitrary outsiders) can join with teacher privileges. As demonstrated in Chen et al.’s work, blockchain-enforced RBAC for video data prevented unauthorized access effectively

link.springer.com

.

Sovereign Digital Identity for Education

In the context of academic/government systems, digital sovereignty often implies that users/institutions fully control their data and identities. Researchers are exploring self-sovereign identity (SSI) on blockchain for education. Chan et al. (2025) survey “blockchain-assisted self-sovereign identities” (Ba-SSI) in education, arguing that students and faculty should own their credentials on a decentralized ledger

mdpi.com

. This model shifts identity management from central servers to distributed ledgers: individuals keep private keys and authorize sharing of their data without trusting any single authority. As the survey notes, Ba-SSI can “ensure security, interoperability, and scalability” in educational identity systems

mdpi.com

. A sovereign meeting platform might leverage SSI: users authenticate via blockchain-held credentials (or verifiable credentials from a permissioned blockchain like Hyperledger Indy), ensuring only authorized academic members participate. This aligns with digital sovereignty goals by avoiding dependence on foreign or commercial identity providers and giving institutions direct control over their network.

Key References: The above concepts are grounded in recent peer-reviewed work. Notably, Yilmaz et al. (IEEE ISNCC 2020) propose a blockchain-anchored identity scheme for WebRTC

researchgate.net

researchgate.net

. Mansell et al. (Algorithms 2021) detail a blockchain audit-trail design

mdpi.com

mdpi.com

. Chen et al. (Cluster Computing 2024) demonstrate Hyperledger Fabric with hybrid encryption and RBAC for secure video data

link.springer.com

link.springer.com

. And Chan et al. (Blockchains 2025) survey blockchain-based self-sovereign identity in education

mdpi.com

. These works provide a strong technical foundation for building a secure, auditable, and sovereign online meeting system for academic and government use.

A Secure and Personalized Online Meeting System for AICTE

<p>Abhishek Muthanna K Department of Computer Science and Engineering Presidency University Bengaluru, India Abhishek.20221CSE0110@presidencyuniversity.in</p>	<p>Chilmakuri Varun Department of Computer Science and Engineering Presidency University Bengaluru, India Chilmakuri.20221CSE0116@presidencyuniversity.in</p>
<p>Amara Hema Harshith Department of Computer Science and Engineering Presidency University Bengaluru, India AMARA.20221CSE0010@presidencyuniversity.in</p>	<p>Dr. Hasan Hussain S Department of Computer Science and Engineering Presidency University Bengaluru, India hasan.hussain@presidencyuniversity.in</p>

Abstract— This research paper describes the conceptualization, engineering, and deployment of a safe personalized online meeting system for the All India Council for Technical Education (AICTE). As AICTE engages in numerous high-stakes meetings with universities, faculty, and Ministry representatives, the study fills an important resource gap for a trustful system with robust data confidentiality, integrity, and ultimately verifiable auditability. The system architecture uses real-time audio/video communication via web RTC (Web Real-Time Communication) and a scalable self-hosted backend system. The primary innovation consists of combining a private permissioned blockchain module (Hyperledger Fabric) as a layer on top of the backend to provide an immutable, tamper-proof, and cryptographically verifiable audit log for all critical meeting activity (e.g. user authentication, file sharing, accessing meeting recordings). The system architecture provides enhanced security, verifiable non-repudiation, and a deeply personalized role-based access control (RBAC) system mapped directly to AICTE's organizational structure. The system addresses the priorities of Digital India and the UN Sustainable Development Goal (SDG) on Industry, Innovation and Infrastructure by enabling secure, resilient, and sovereign communication infrastructures for public sector education governance.

Index Terms—Online Meeting, AICTE, WebRTC, Blockchain, Hyperledger Fabric, Cybersecurity, Secure Communication, E-Governance, Verifiable Audit, Data Sovereignty, Digital India, SDG 9.

I. INTRODUCTION

The digital transformation of governance and public administration has accelerated the transition to online meeting and collaboration platforms [1]. For a national regulatory and accreditation body, like the All India Council for Technical Education (AICTE), hosting meetings remotely is no longer a luxury but a need. AICTE is engaged with thousands of technical institutions, faculty, industry stakeholders, and Ministry officials and meeting activities engage from regular meeting to sensitive committee discussions, policy makers review, and grievance hearings [2].

However, reliance on moving communications to digital channels has significant attack surfaces. AICTE uses commercial off-the-shelf (COTS) video communications systems, hosted often on public clouds based outside national jurisdiction. This poses a number of critical challenges around data privacy, surveillance, unauthorized access, and data sovereignty [4]. In this context, a security event could easily compromise sensitive national policy documents, institutional

data, or personal identifiers of stakeholders.

A. Problem Statement and Motivation

AICTE meetings regularly have discussions and sharing of sensitive, confidential, and "For Official Use Only" files. Depending on a third-party platform exposes an unacceptable risk. Standard platforms can offer audit logs; however, audit logs are typically opaque, vendor-controlled, and, most importantly, mutable. There are no guarantees that the logs have not been modified by an administrator or a savvy attacker to conceal the breach. The absence of verifiable non-repudiation is a fundamental flaw for e-governance high security applications.

Furthermore, COTS platforms are designed with generic access control models (host, co-host, participant) which do not map to the complex organizational hierarchy of a government agency like AICTE. A "one-size-fits-all" security model is insufficient where permissions must be defined more granularly based on roles such as, "Institute Head," "Faculty," "HOD," "Ministry Official," or "Committee Member."

This research is driven by the critical need for tailored, self-hosted, and highly secure meeting software that satisfies three essential requirements:

- 1) **Data Sovereignty:** The platform must be hosted either on AICTE's own data center or on a trusted national cloud environment. This implies that no data (streams, files, metadata, etc.) is allowed to leave the jurisdiction of the data center or the national cloud.
- 2) **Personalized Security:** The system must include a deep Role-Based Access Control (RBAC) model that reflects the specific organizational structure of AICTE.
- 3) **Verifiable Auditability:** The system must provide a non-repudiable, tamper-proof audit trail for all critical actions, which can be cryptographically verified by any authorized party.

B. Alignment with National Goals and SDGs

This project directly supports the vision of a self-reliant "Digital India" by creating secure, sovereign digital infrastructure for critical state functions [9]. It also contributes to UN Sustainable Development Goal 9 (Industry, Innovation, and Infrastructure) by "building resilient infrastructure, promoting

PUBLICATION

Dear CHILMAKURI VARUN,

Thanks for considering IJRPR. Your paper is successfully received on **Date 02/12/2025 10:09 PM** via our Online Portal. Your **Paper ID** is **IJRPR-191405**.

Paper Title : A Secure and Personalized Online Meeting System for AICTE

You can Check the Status of paper by using above Paper ID [Check Paper Status](#)

The Review result of your paper will be mailed to you once the review process is completed. For any future communication you are advised to refer your **Paper ID IJRPR-191405**.

With Warm Regards

Editor-In Chief

International Journal of Research Publication and Reviews (IJRPR)

<http://www.ijrpr.com>

Dear CHILMAKURI VARUN,

This is to inform you that your paper has been accepted for publication in IJRPR, which will be published in current issue. The Unique ID of your paper is IJRPR-191405.

Paper Title: A Secure and Personalized Online Meeting System for AICTE

You are advised to complete the process for the publication of your research paper.

1. Fill & Send Copyright Transfer Form (download from our website <https://www.ijrpr.com/download/COPYRIGHT-FORM.pdf>)

(Note: Take a print of copyright form, fill it, scan it & send us via mail to contactusijrpr@gmail.com or send it in word format with digital/scanned signatures. Do not forget to mention your paper ID in subject of mail.)

Alternatively you can submit the Online Copyright Form by clicking the following link

<https://forms.gle/5PcfsA3eqLkhR3c6>

2. Submit your Publication fee

(Note: Send the Payment receipt in image/word/pdf format along with copyright form to contactusijrpr@gmail.com Do not forget to mention your paper ID in subject of mail.)

3. Paper will be published within 24 to 36 working Hours after the completion of step 1 & 2. We will provide the Soft Individual copy of the certificates to all authors of a research paper.

Do not forget to mention your paper ID in subject.

Publication Fee:

International Authors	17 US Dollars (Click Here to Pay)
Indian Authors (Author Affiliation in Paper Must be in Indian Territory)	Rs. 599 (Click Here to Pay)
E-Certificate	Free

Publication Charge includes:

- Publication of one entire research paperOnline
- Individual Soft Copy of Certificate to all author of paper.
- Editorial Fee
- Indexing, maintenance of link resolvers and journal infrastructures.

SIMILARITY CHECK

Appendix

Appendix A : Project Images

This appendix contains screenshots of the developed system to demonstrate its functional behavior in real scenarios.

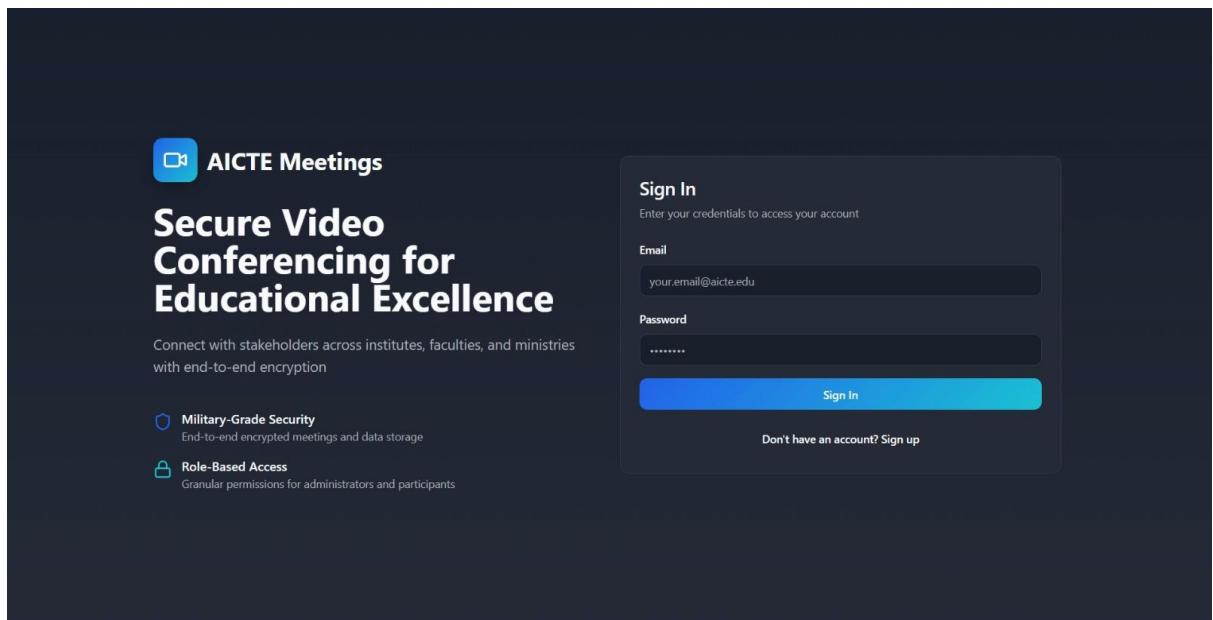


Fig A : Sign in page

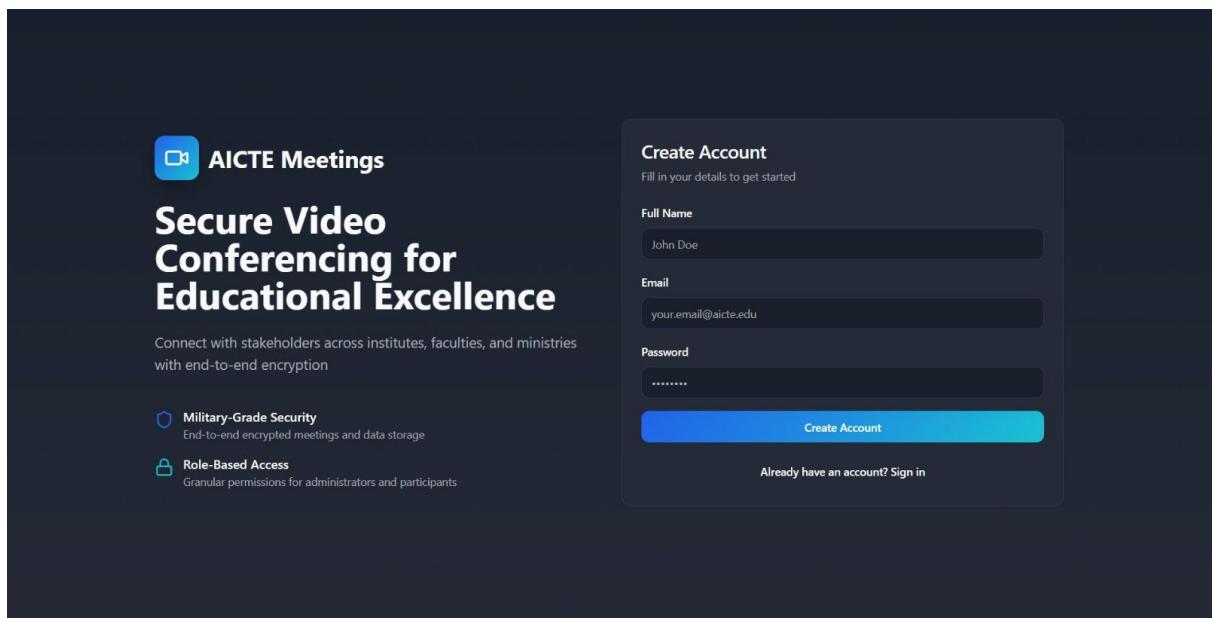


Fig B : Sign up page

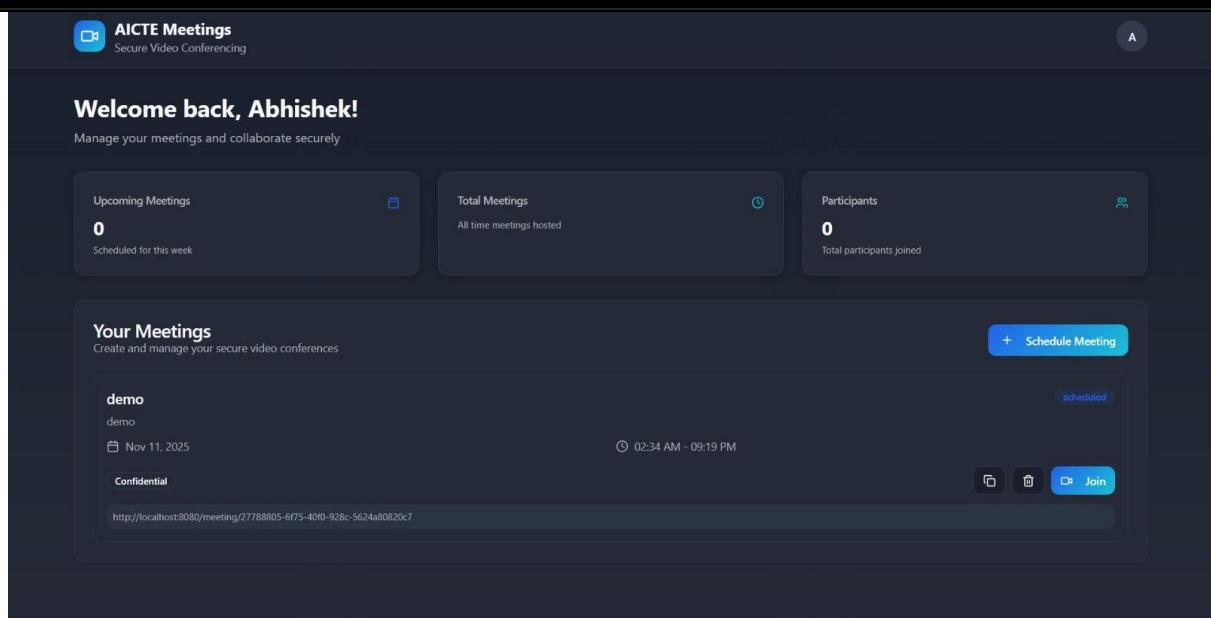


Fig C : Home Screen

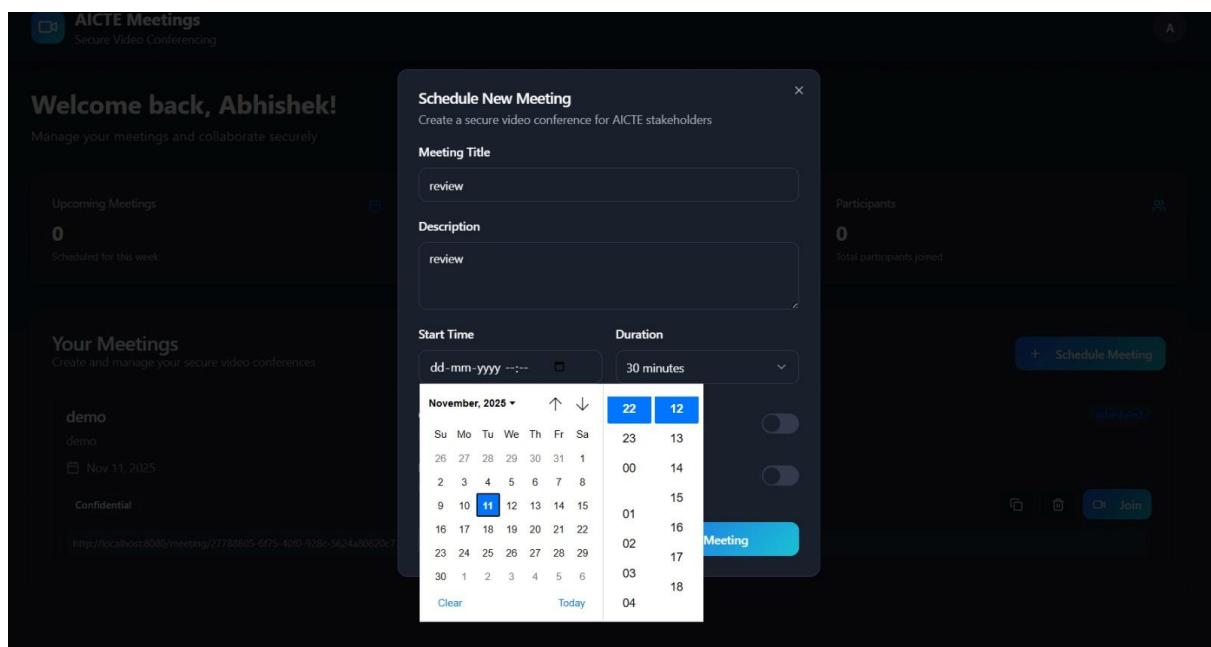


Fig D : Meeting Scheduler

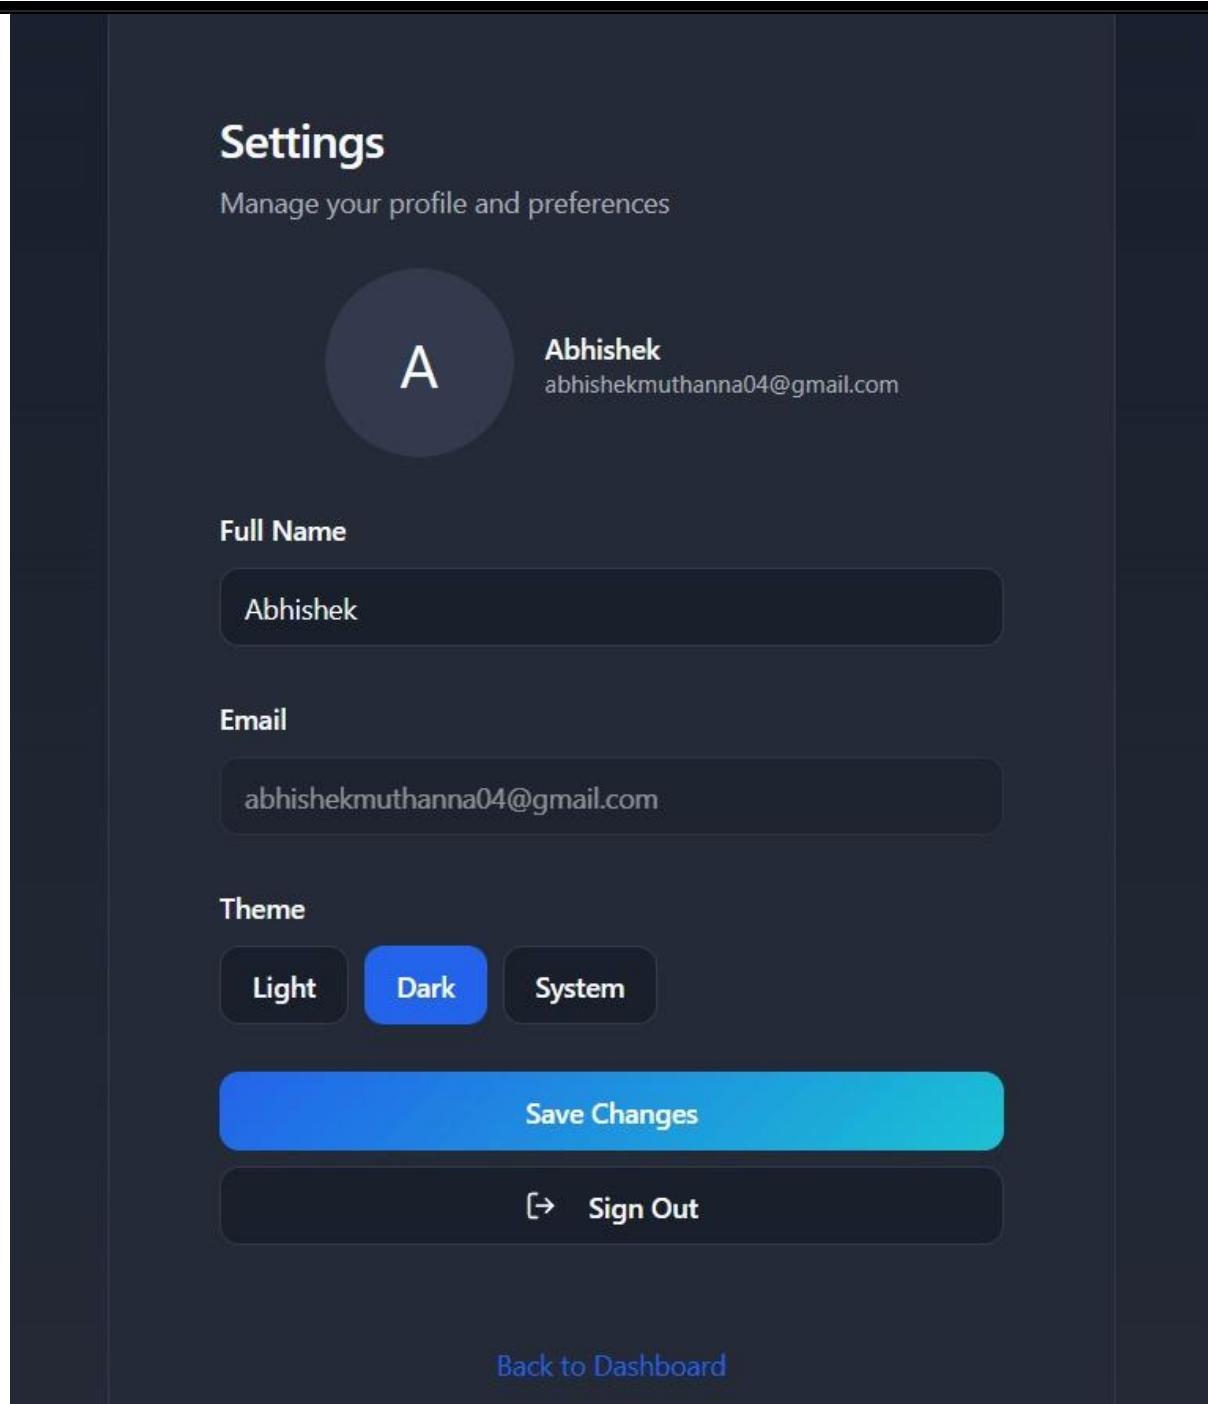


Fig E : Profile Settings



Fig F : Main Menu

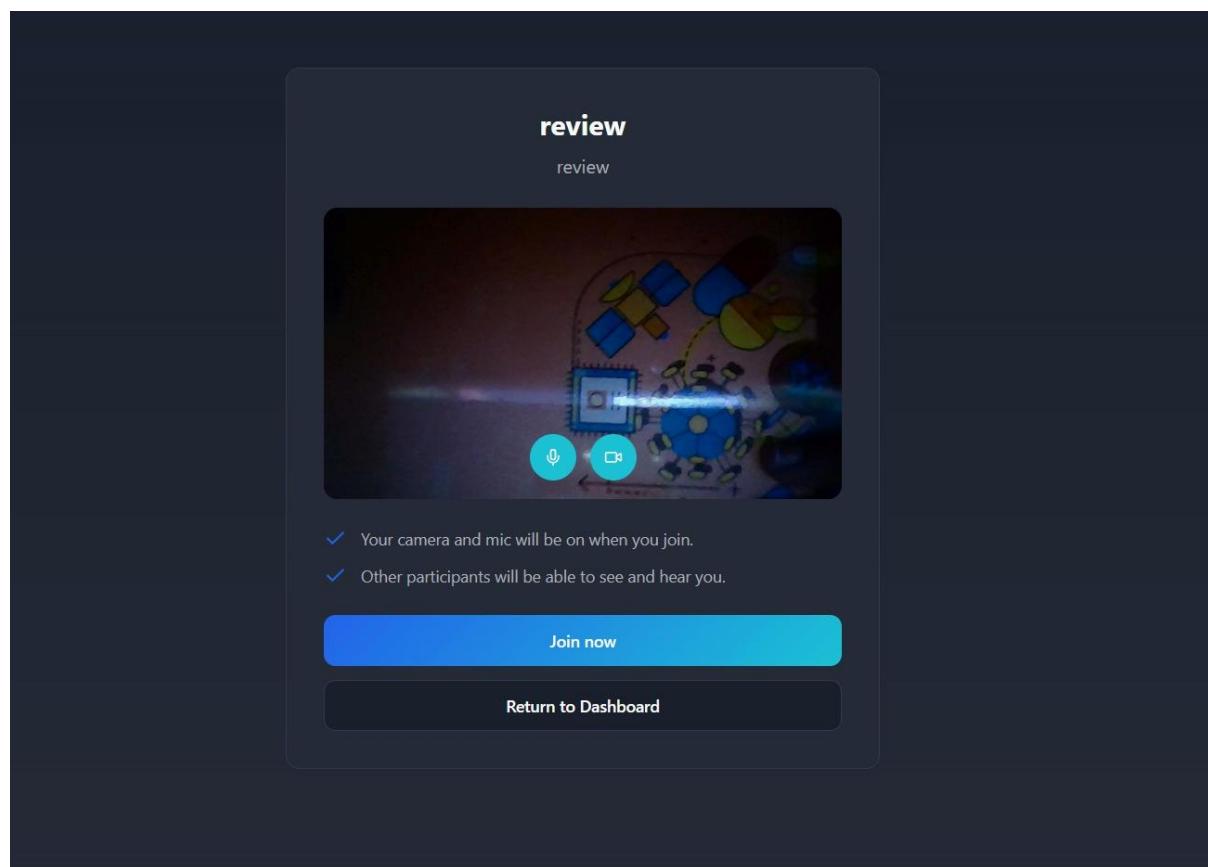


Fig G : Participant waiting page for meeting.

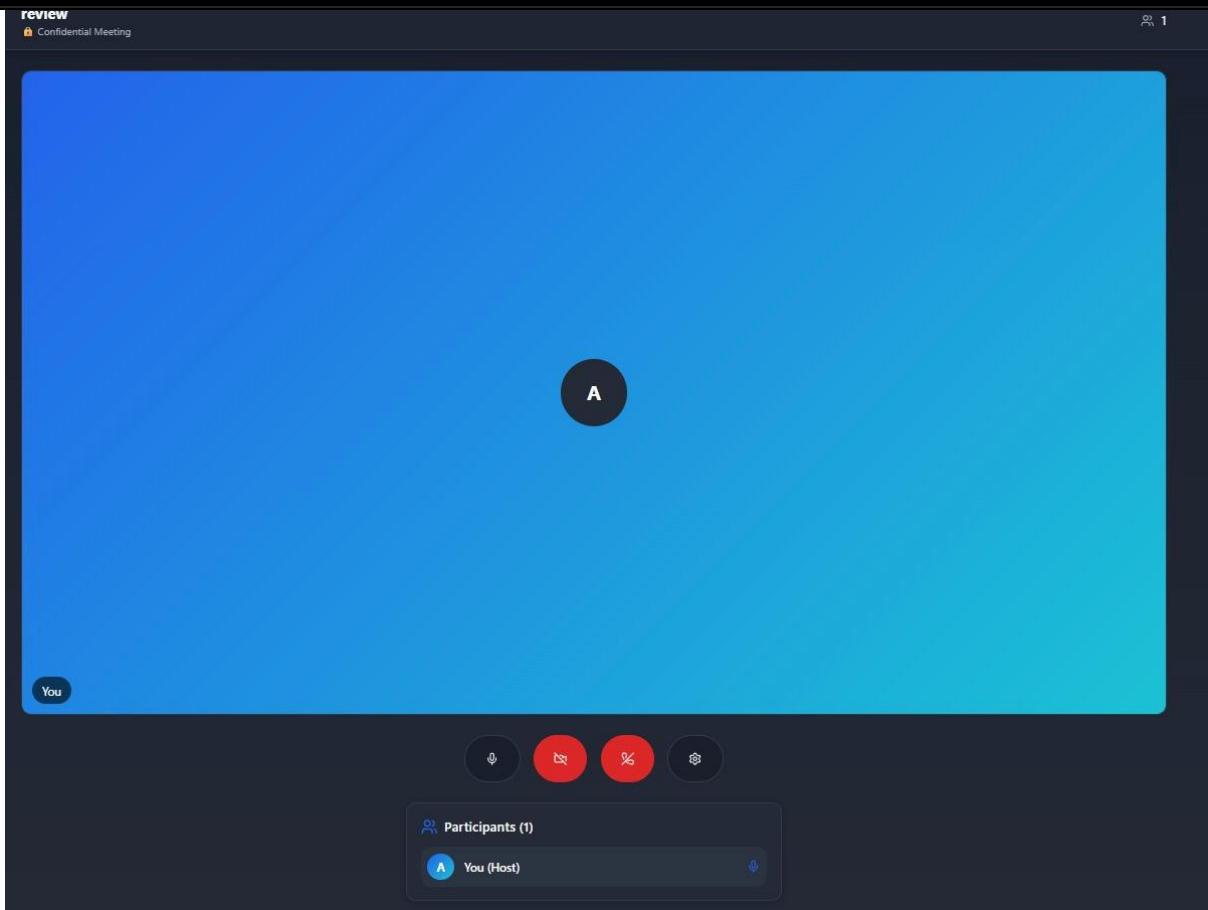


Fig H : Meeting platform.

Appendix B: Code Snippets

Deterministic Initiator Logic

```

import { hmac } from "utils/crypto"; // Hypothetical crypto util

/**
 * Deterministically decides which peer should initiate the WebRTC offer.
 * Implements logic defined in Methodology Section 3.3.1.
 * @param localId - UUID of the local user
 * @param remoteId - UUID of the remote peer
 * @param meetingEpoch - Meeting start timestamp (ISO string)
 * @returns boolean - True if local user should initiate offer
 */
function decideInitiator(localId: string, remoteId: string, meetingEpoch: string):
    // 1. Prevent self-connection
    if (localId === remoteId) return false;

    // 2. Generate deterministic hash for the pair
    // Order ensures hash is identical for both peers: min(A,B) + max(A,B)
    const sortedIds = [localId, remoteId].sort();
    const pairString = `${meetingEpoch}:${sortedIds[0]}:${sortedIds[1]}`;

    // 3. Compute Hash (HMAC-SHA256)
    const pairHash = hmac(pairString, process.env.SYSTEM_SECRET);

    // 4. Extract Seed (first 8 bytes as integer)
    const seed = parseInt(pairHash.substring(0, 8), 16);

    // 5. Epoch-based Randomized Tie-Breaker
    if (seed % 2 === 0) {
        // Lexicographic Ascending
        return localId < remoteId;
    } else {
        // Lexicographic Descending
        return localId > remoteId;
}

```

Fig B1 : Deterministic Initiator Selection Algorithm (TypeScript).

Hyperledger Fabric Chaincode (Audit Logging)

```

import (
    "encoding/json"
    "fmt"
    "github.com/hyperledger/fabric-contract-api-go/contractapi"
)

// AuditEvent defines the structure for anchored logs
type AuditEvent struct {
    EventID string `json:"eventId"`
    MeetingID string `json:"meetingId"`
    ActorID string `json:"actorId"`
    DataHash string `json:"dataHash"` // SHA-256 of the off-chain payload
    Timestamp string `json:"timestamp"`
}

// AuditContract provides functions for managing audit logs
type AuditContract struct {
    contractapi.Contract
}

// LogEvent anchors a new audit hash to the ledger
// Corresponds to design specification in Section 3.4
func (c *AuditContract) LogEvent(ctx contractapi.TransactionContextInterface, eventID string) error {
    // 1. Validate Identity (MSP Check)
    clientMSP, err := ctx.GetClientIdentity().GetMSPID()
    if err != nil {
        return fmt.Errorf("failed to get client MSP: %v", err)
    }
    if clientMSP != "AICTE_Org1MSP" {
        return fmt.Errorf("unauthorized organization")
    }

    // 2. Check for Duplicates
    exists, err := ctx.GetStub().GetState("EVT:" + eventID)
    if err != nil {
        return err
    }
    if exists != nil {
        return fmt.Errorf("event %s already exists", eventID)
    }

    // 3. Create Event Object
    event := AuditEvent{
        EventID: eventID,
        MeetingID: meetingID,
        ActorID: actorID,
        DataHash: dataHash,
        Timestamp: timestamp,
    }

    // 4. Serialize and Write to Ledger
    eventJSON, err := json.Marshal(event)
    if err != nil {
        return err
    }

    return ctx.GetStub().PutState("EVT:"+eventID, eventJSON)
}

```

Fig B2 : Hyperledger Fabric Chaincode for Immutable Event Logging (Go)

Supabase Signaling Logic

```

interface SignalingMessage {
    type: "OFFER" | "ANSWER" | "ICE" | "EVENT";
    from: string; // User UUID
    to: string; // Target UUID or "ALL"
    seq: number; // Monotonic sequence number for ordering
    meetingId: string; // Context verification
    timestamp: string; // ISO8601 UTC
    payload: RTCSessionDescriptionInit | RTCIceCandidateInit | any;
    signature?: string; // HMAC for integrity
}

// Function to subscribe to meeting signaling channel
async function subscribeToSignaling(meetingId: string, supabase: SupabaseClient)
    const channelName = `meeting:${{meetingId}}:control`;

    const channel = supabase.channel(channelName)
        .on('broadcast', { event: 'signal' }, (payload) => {
            const msg = payload.payload as SignalingMessage;

            // 1. Validate Message Context
            if (msg.meetingId !== currentMeetingId) return;

            // 2. Route to WebRTC Manager
            handleIncomingSignal(msg);
        })
        .subscribe((status) => {
            if (status === 'SUBSCRIBED') {
                console.log(`Joined signaling channel: ${channelName}`);
            }
        });
}

return channel;

```

Fig B3 : Real-time Signaling Envelope Interface (TypeScript)

RBAC Policy Evaluation

```

const policyStore = {
    "MEETING_RECORDING": {
        "START": {
            allowedRoles: ["HOST", "CO_HOST", "MINISTRY_OFFICIAL"],
            conditions: ["is_meeting_active", "storage_quota_available"]
        },
        "STOP": {
            allowedRoles: ["HOST", "CO_HOST"]
        }
    },
    "FILE_SHARE": {
        "UPLOAD": {
            allowedRoles: ["HOST", "FACULTY", "STUDENT"],
            conditions: ["file_type_allowed", "virus_scan_pass"]
        }
    }
};

function evaluatePolicy(actorRole, action, resourceType, context) {
    // 1. Retrieve Policy
    const policy = policyStore[resourceType]?.[action];
    if (!policy) return false; // Default Deny

    // 2. Check Role (Direct Allow)
    if (!policy.allowedRoles.includes(actorRole)) {
        console.warn(`Access Denied: Role ${actorRole} cannot perform ${action}`);
        return false;
    }

    // 3. Evaluate Contextual Conditions
    if (policy.conditions) {
        for (const condition of policy.conditions) {
            if (!checkCondition(condition, context)) {
                return false;
            }
        }
    }

    return true; // Access Granted
}

```

Fig B4 : Role-Based Access Control (RBAC) Policy Evaluator

Appendix C: Git Hub

The screenshot shows a GitHub repository page for the project 'major_project-'. The repository is public and has 1 branch and 0 tags. The main file list shows numerous files uploaded by 'AbhishekMuthannaK' via upload, all marked as 'now'. The files include .env, .gitignore, .package-lock.json, README.md, bun.lockb, components.json, config.json, eslint.config.js, index.html, package-lock.json, package.json, postcss.config.js, readme.md, sponsors.md, tailwind.config.ts, tsconfig.app.json, tsconfig.json, tsconfig.node.json, and vite.config.ts. The repository has 2 commits, 0 stars, 0 forks, and 0 watching. It also lists 'About', 'Readme', 'Activity', 'Releases', and 'Packages' sections.

File	Uploader	Action	Last Commit
.env	AbhishekMuthannaK	Add files via upload	now
.gitignore	AbhishekMuthannaK	Add files via upload	now
.package-lock.json	AbhishekMuthannaK	Add files via upload	now
README.md	AbhishekMuthannaK	Add files via upload	now
bun.lockb	AbhishekMuthannaK	Add files via upload	now
components.json	AbhishekMuthannaK	Add files via upload	now
config.json	AbhishekMuthannaK	Add files via upload	now
eslint.config.js	AbhishekMuthannaK	Add files via upload	now
index.html	AbhishekMuthannaK	Add files via upload	now
package-lock.json	AbhishekMuthannaK	Add files via upload	now
package.json	AbhishekMuthannaK	Add files via upload	now
postcss.config.js	AbhishekMuthannaK	Add files via upload	now
readme.md	AbhishekMuthannaK	Add files via upload	now
sponsors.md	AbhishekMuthannaK	Add files via upload	now
tailwind.config.ts	AbhishekMuthannaK	Add files via upload	now
tsconfig.app.json	AbhishekMuthannaK	Add files via upload	now
tsconfig.json	AbhishekMuthannaK	Add files via upload	now
tsconfig.node.json	AbhishekMuthannaK	Add files via upload	now
vite.config.ts	AbhishekMuthannaK	Add files via upload	now

Fig C1 : Git Hub Repository

Git Hub Link : https://github.com/AbhishekMuthannaK/major_project-.git